

# ENSEMBLE METHODS

- Ensemble of same classifiers
  - Bagging and random forest
  - Boosting and gradient boosted tree
- Ensemble of different classifiers

# REVIEW: BAGGING

Bootstrap samples



Original dataset

Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
9	2013-11-02 The Hunger Games: Catching Fire	13000000	424885047	Francis Lawrence	PG-13	146
1	2013-08-02 Iron Man 3	120000000	1214953494	Shane Black	PG-13	139
2	2013-11-02 Frozen	10000000	401738039	Chris Buck/Jennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	369811265	Pierre Coffin/Chris Renaud	PG	98
4	2013-08-14 Man of Steel	220000000	219145118	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	10000000	174302705	Alfonso Cuarón	PG-13	91
6	2013-08-21 Monsters University	N/A	284482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	N/A	256388650	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238779550	Justin Lin	PG-13	130
10	2013-05-08 On the Beach and Powerful	21000000	224917625	Sam Raimi	PG	127
11	2013-05-16 Star Trek Into Darkness	99000000	228779861	J.J. Abrams	PG-13	123
12	2013-11-08 Thor: The Dark World	170000000	206302140	Alan Taylor	PG-13	120
13	2013-08-01 World War Z	160000000	203307111	Mark Forster	PG-13	116
14	2013-05-03 The Conjuring	13000000	187186825	Kris Zschock/Chris Sanders	PG	98
14	2013-08-28 The Heat	4300000	158502188	Paul Feig	R	117
16	2013-08-07 When We Were Millions	3700000	103294119	Ramsey Marshall Thudler	R	110
16	2013-12-13 American Hustle	4000000	102177907	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	10000000	144840419	Baz Luhrmann	PG-13	143

Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
9	2013-11-02 The Hunger Games: Catching Fire	13000000	424885047	Francis Lawrence	PG-13	146
1	2013-08-02 Iron Man 3	120000000	1214953494	Shane Black	PG-13	139
2	2013-11-02 Frozen	10000000	401738039	Chris Buck/Jennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	369811265	Pierre Coffin/Chris Renaud	PG	98
4	2013-08-14 Man of Steel	220000000	219145118	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	10000000	174302705	Alfonso Cuarón	PG-13	91
6	2013-08-21 Monsters University	N/A	284482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	N/A	256388650	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238779550	Justin Lin	PG-13	130
10	2013-05-08 On the Beach and Powerful	21000000	224917625	Sam Raimi	PG	127
11	2013-05-16 Star Trek Into Darkness	99000000	228779861	J.J. Abrams	PG-13	123
12	2013-11-08 Thor: The Dark World	170000000	206302140	Alan Taylor	PG-13	120
13	2013-08-01 World War Z	160000000	203307111	Mark Forster	PG-13	116
14	2013-05-03 The Conjuring	13000000	187186825	Kris Zschock/Chris Sanders	PG	98
14	2013-08-28 The Heat	4300000	158502188	Paul Feig	R	117
16	2013-08-07 When We Were Millions	3700000	103294119	Ramsey Marshall Thudler	R	110
16	2013-12-13 American Hustle	4000000	102177907	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	10000000	144840419	Baz Luhrmann	PG-13	143

Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
9	2013-11-02 The Hunger Games: Catching Fire	13000000	424885047	Francis Lawrence	PG-13	146
1	2013-08-02 Iron Man 3	120000000	1214953494	Shane Black	PG-13	139
2	2013-11-02 Frozen	10000000	401738039	Chris Buck/Jennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	369811265	Pierre Coffin/Chris Renaud	PG	98
4	2013-08-14 Man of Steel	220000000	219145118	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	10000000	174302705	Alfonso Cuarón	PG-13	91
6	2013-08-21 Monsters University	N/A	284482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	N/A	256388650	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238779550	Justin Lin	PG-13	130
10	2013-05-08 On the Beach and Powerful	21000000	224917625	Sam Raimi	PG	127
11	2013-05-16 Star Trek Into Darkness	99000000	228779861	J.J. Abrams	PG-13	123
12	2013-11-08 Thor: The Dark World	170000000	206302140	Alan Taylor	PG-13	120
13	2013-08-01 World War Z	160000000	203307111	Mark Forster	PG-13	116
14	2013-05-03 The Conjuring	13000000	187186825	Kris Zschock/Chris Sanders	PG	98
14	2013-08-28 The Heat	4300000	158502188	Paul Feig	R	117
16	2013-08-07 When We Were Millions	3700000	103294119	Ramsey Marshall Thudler	R	110
16	2013-12-13 American Hustle	4000000	102177907	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	10000000	144840419	Baz Luhrmann	PG-13	143

Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
9	2013-11-02 The Hunger Games: Catching Fire	13000000	424885047	Francis Lawrence	PG-13	146
1	2013-08-02 Iron Man 3	120000000	1214953494	Shane Black	PG-13	139
2	2013-11-02 Frozen	10000000	401738039	Chris Buck/Jennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	369811265	Pierre Coffin/Chris Renaud	PG	98
4	2013-08-14 Man of Steel	220000000	219145118	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	10000000	174302705	Alfonso Cuarón	PG-13	91
6	2013-08-21 Monsters University	N/A	284482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	N/A	256388650	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238779550	Justin Lin	PG-13	130
10	2013-05-08 On the Beach and Powerful	21000000	224917625	Sam Raimi	PG	127
11	2013-05-16 Star Trek Into Darkness	99000000	228779861	J.J. Abrams	PG-13	123
12	2013-11-08 Thor: The Dark World	170000000	206302140	Alan Taylor	PG-13	120
13	2013-08-01 World War Z	160000000	203307111	Mark Forster	PG-13	116
14	2013-05-03 The Conjuring	13000000	187186825	Kris Zschock/Chris Sanders	PG	98
14	2013-08-28 The Heat	4300000	158502188	Paul Feig	R	117
16	2013-08-07 When We Were Millions	3700000	103294119	Ramsey Marshall Thudler	R	110
16	2013-12-13 American Hustle	4000000	102177907	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	10000000	144840419	Baz Luhrmann	PG-13	143



# REVIEW: RANDOM FORESTS

- Bagged classifier using decision trees
  - Each split only considers a **random group of features**
  - **Tree is grown to maximum size without pruning**
  - Final predictions obtained by aggregating over the B trees

$$\hat{f}_{\text{rf}}^B(\mathbf{x}) = \frac{1}{B} \sum_b T(\mathbf{x}; \theta_b)$$

- Reduce variance (at the cost of slight increase in bias compared to bagged trees)

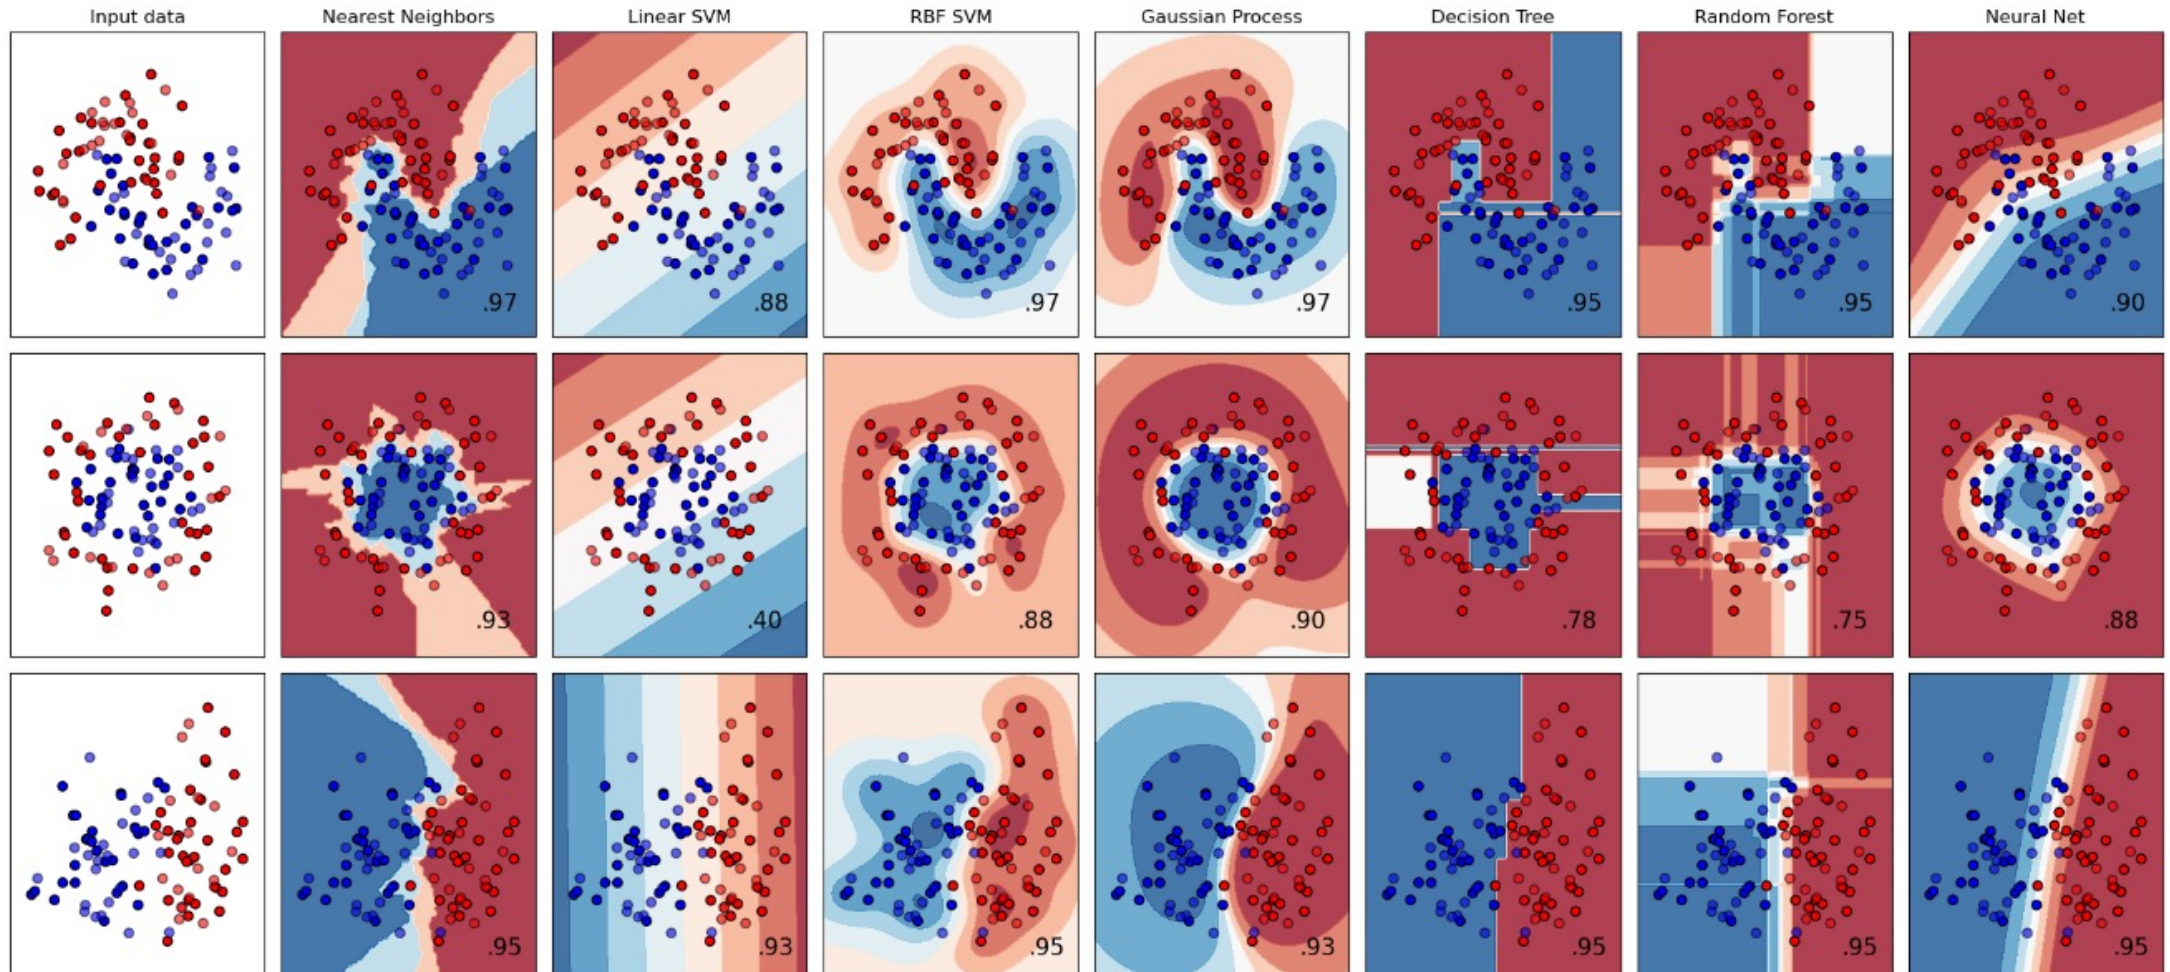
# RANDOM FOREST: ADVANTAGES

- State of the art method, one of the most accurate general-purpose learners available
- Handles a large number of input variables without overfitting (variance reduction)
- Robust to errors and outliers
- Can model non-linear boundaries
- Gives variable importance and out of bag error rates
- Easy to train and tune, easily parallelized by training

# RANDOM FOREST: DISADVANTAGES

- Loss of interpretability (no decision rules)
- Difficult to analyze as an algorithm and mathematical properties still largely unknown
- Large number of trees is memory-intensive
- Bias towards categorical variables with larger number of levels

# RANDOM FOREST



# PREVIEW: HOMEWORK #5

- Almost Random Forest
- Instead of choosing a random subset of features for each split, choose a random subset of features that the tree will be created on (the same subset is used as candidates from all splits)

# SKLEARN: RANDOM FOREST

- `sklearn.ensemble.RandomForestClassifier`
  - `n_estimators`, default=100
  - `max_features`: {"sqrt", "log2", None}, default="sqrt"

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>



# ENSEMBLE

- Ensemble of same classifiers
  - Bagging and random forest
  - Boosting and gradient boosted tree
- Ensemble of different classifiers

# BOOSTING AND GRADIENT BOOSTED TREE

CS 334: Machine Learning



GROUP ACTIVITY

# STUDY FOR AN EXAM

*Exam*

-----

-----

-----

-----

-----

-----

-----

How would you study for an exam  
using a practice exam (given solution keys)?

# STUDY FOR AN EXAM

*Exam*

✓	_____
✗	_____
✓	_____
✓	_____
✓	_____
✗	_____
✗	_____

*Exam*


✓	_____
✓	_____
✓	_____
✓	_____
✓	_____
✓	_____
✗	_____

*Exam*

✓	_____
✓	_____
✓	_____
✓	_____
✓	_____
✓	_____
✓	_____


# BAGGING

- Combine output of many classifiers (full sized trees)
- The learners are combined with equal weights
- Independently train each learner on a bootstrap



# BOOSTING

- Combine output of many **weak** learners (stumps)
- The weak learners are combined with **weights**
- **Sequentially** train weak learners by compensating the **shortcomings of previous learner**



# BOOSTING

- AdaBoost (1997): shortcomings are identified by high weight (misclassified) data points
- Gradient Boosting (1999): shortcomings are identified by residuals or gradients of the loss function (generalize AdaBoost)



# ADABOOST

- Maintain a weight distribution over data points (importance of having the data point correctly classified)
- The weights are initialized as uniform and updated after each classifier is trained to guide next classifier (correctly classified examples have weight decreased and incorrect examples increased)
- Final classifier is weighted vote of individual classifiers, based on weighted error of each classifier

# ADABOOST

---

**Algorithm 32**  $\text{ADABOOST}(\mathcal{W}, \mathcal{D}, K)$ 

---

```
1:  $d^{(0)} \leftarrow \langle \frac{1}{N}, \frac{1}{N}, \dots, \frac{1}{N} \rangle$            // Initialize uniform importance to each example
2: for  $k = 1 \dots K$  do
3:    $f^{(k)} \leftarrow \mathcal{W}(\mathcal{D}, d^{(k-1)})$            // Train  $k$ th classifier on weighted data
4:    $\hat{y}_n \leftarrow f^{(k)}(x_n), \forall n$                // Make predictions on training data
5:    $\hat{\epsilon}^{(k)} \leftarrow \sum_n d_n^{(k-1)} [y_n \neq \hat{y}_n]$  // Compute weighted training error
6:    $\alpha^{(k)} \leftarrow \frac{1}{2} \log \left( \frac{1 - \hat{\epsilon}^{(k)}}{\hat{\epsilon}^{(k)}} \right)$  // Compute “adaptive” parameter
7:    $d_n^{(k)} \leftarrow \frac{1}{Z} d_n^{(k-1)} \exp[-\alpha^{(k)} y_n \hat{y}_n], \forall n$  // Re-weight examples and normalize
8: end for
9: return  $f(\hat{x}) = \text{sgn} [\sum_k \alpha^{(k)} f^{(k)}(\hat{x})]$  // Return (weighted) voted classifier
```

---

Weighted entropy/gini  
or weighted sampling

Assuming error  $< 0.5$   
Smaller error, great alpha

Decrease weight for  
correct samples, increase  
weight for incorrect  
samples

Smaller error  $\rightarrow$  more significant re-weighting, why?

# ADABOOST





Figure: AdaBoost. Source: Figure 1.1 of [Schapire and Freund, 2012]

# ADABOOST

$$H = \text{sign} \left( 0.42 \begin{array}{|c|} \hline \text{dark gray} \\ \hline \end{array} + 0.65 \begin{array}{|c|} \hline \text{dark gray} \\ \hline \end{array} + 0.92 \begin{array}{|c|} \hline \text{dark gray} \\ \hline \end{array} \right)$$



# BOOSTING

- AdaBoost: shortcomings are identified by high weight (misclassified) data points
- Gradient Boosting: shortcomings are identified by residuals or gradients of the loss function (generalize AdaBoost)

# GRADIENT BOOSTING

- Gradient descent + boosting
- Powerful algorithm that can be used for regression, classification, ranking
- Data mining competition winner most likely uses this algorithm



GROUP ACTIVITY

# ADDITIVE MODEL

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

- We have training data,  $(x_1, y_1), \dots, (x_n, y_n)$ , and task to fit model  $f(x)$  for  $y$  to minimize square loss
- A friend gives you a model  $f$ , with some mistakes
- You are not allowed to remove anything from  $f$  or change any parameter in  $f$
- You are allowed to add additional models  $h$  to  $f$ , so new prediction is  $f(x) + h(x)$

What would you do?



# ADDITIVE MODEL

- Simple solution:


Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2

$$f(x_1) + h(x_1) = y_1$$

$$f(x_2) + h(x_2) = y_2$$

$$\vdots$$

$$f(x_n) + h(x_n) = y_n$$

$$\vdots$$


$$h(x_1) = y_1 - f(x_1)$$

$$h(x_2) = y_2 - f(x_2)$$

$$\vdots$$

$$h(x_n) = y_n - f(x_n)$$

$$\vdots$$

We can train a regression tree for residual  $h$ ,  
then add the tree to original model  $f$

# GRADIENT BOOSTING: REGRESSION

- Initialize a constant model  $f(x)$  for data  $(x_1, y_1), \dots, (x_n, y_n)$
- Fit a regression tree  $h(x)$  to new data,  $(x_1, y_1 - f(x_1)), \dots, (x_n, y_n - f(x_n))$
- $y_i - f(x_i)$  are **pseudo residuals** - parts that existing model  $f$  cannot do well
- Update  $f(x) + h(x)$  as prediction, if it is not satisfactory, reiterate

# GRADIENT BOOSTED TREE

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

# GRADIENT BOOSTED TREE

71.2

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

# GRADIENT BOOSTED TREE

71.2

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2

# GRADIENT BOOSTED TREE

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2

71.2




# GRADIENT BOOSTED TREE

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2


71.2

+



# GRADIENT BOOSTED TREE

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2





# GRADIENT BOOSTING: REGRESSION

- Initialize a constant model  $f(x)$  for data  $(x_1, y_1), \dots, (x_n, y_n)$
- Fit a regression tree  $h(x)$  to new data,  $(x_1, y_1 - f(x_1)), \dots, (x_n, y_n - f(x_n))$
- $y_i - f(x_i)$  are **pseudo residuals** - parts that existing model  $f$  cannot do well
- Update  $f(x) + h(x)$  as prediction, if it is not satisfactory, reiterate

How does this relate to gradient descent?

# GRADIENT BOOSTING: REGRESSION

- Loss function:  $L(\mathbf{y}, f(\mathbf{X})) = \sum_i \frac{1}{2} (y_i - f(\mathbf{x}_i))^2$
- Treat  $f(\mathbf{x}_i)$  as parameters and take derivatives


$$\frac{\partial L}{\partial f(\mathbf{x}_i)} = f(\mathbf{x}_i) - y_i$$

- Interpret residuals as negative gradients (learning rate = 1 here)

$$f(\mathbf{x}_i) := f(\mathbf{x}_i) + \boxed{y_i - f(\mathbf{x}_i)}$$

← Residual estimated by next tree

$$f(\mathbf{x}_i) := f(\mathbf{x}_i) - 1 \frac{\partial L}{\partial \mathbf{x}_i}$$



# GRADIENT TREE BOOSTING: ALGORITHM

---

**Algorithm 10.3** *Gradient Tree Boosting Algorithm.*

---

1. Initialize  $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$ . ← Constant model

2. For  $m = 1$  to  $M$ :

(a) For  $i = 1, 2, \dots, N$  compute

$$r_{im} = - \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}$$

← Residual of  $i$  or negative gradient

(b) Fit a regression tree to the targets  $r_{im}$  giving terminal regions  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$ .  
←  $J_m$  is the size of the tree


(c) For  $j = 1, 2, \dots, J_m$  compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma)$$

← Average of residuals in each leaf

(d) Update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$ . ← update

3. Output  $\hat{f}(x) = f_M(x)$ .



# SHRINKAGE

---

**Algorithm 10.3** *Gradient Tree Boosting Algorithm.*

---

1. Initialize  $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$ .

2. For  $m = 1$  to  $M$ :

(a) For  $i = 1, 2, \dots, N$  compute

$$r_{im} = - \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}.$$

(b) Fit a regression tree to the targets  $r_{im}$  giving terminal regions  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$ .


(c) For  $j = 1, 2, \dots, J_m$  compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma).$$

(d) Update  $f_m(x) = f_{m-1}(x) + \boxed{\nu} \sum_{j=1}^J \gamma_{jm} I(x \in R_{jm})$ . ←


3. Output  $\hat{f}(x) = f_M(x)$ .

---



$\nu$ : add a shrinkage parameter to control the learning rate of the boosting procedure.


# GRADIENT BOOSTING: ITERATIONS



What do we observe here?

# GRADIENT BOOSTING: ITERATIONS

- Boosting (bias reduction) vs random forest (variance reduction)
- Boosting is additive, reduces bias, so can overfit (variance can increase)  
(like over studying practice exam)
- Use standard mechanism to tune the parameters



# SIZE OF TREES

- $J_m$ : the size of tree for each iteration
- Simplest strategy is to use the same  $J_m = J$
- Best method is to grow small trees with no pruning
- Right size will depend on level of interaction between variables
- ~4-8 leaves works well








Figure 10.9 (Hastie et al.)

# SHRINKAGE

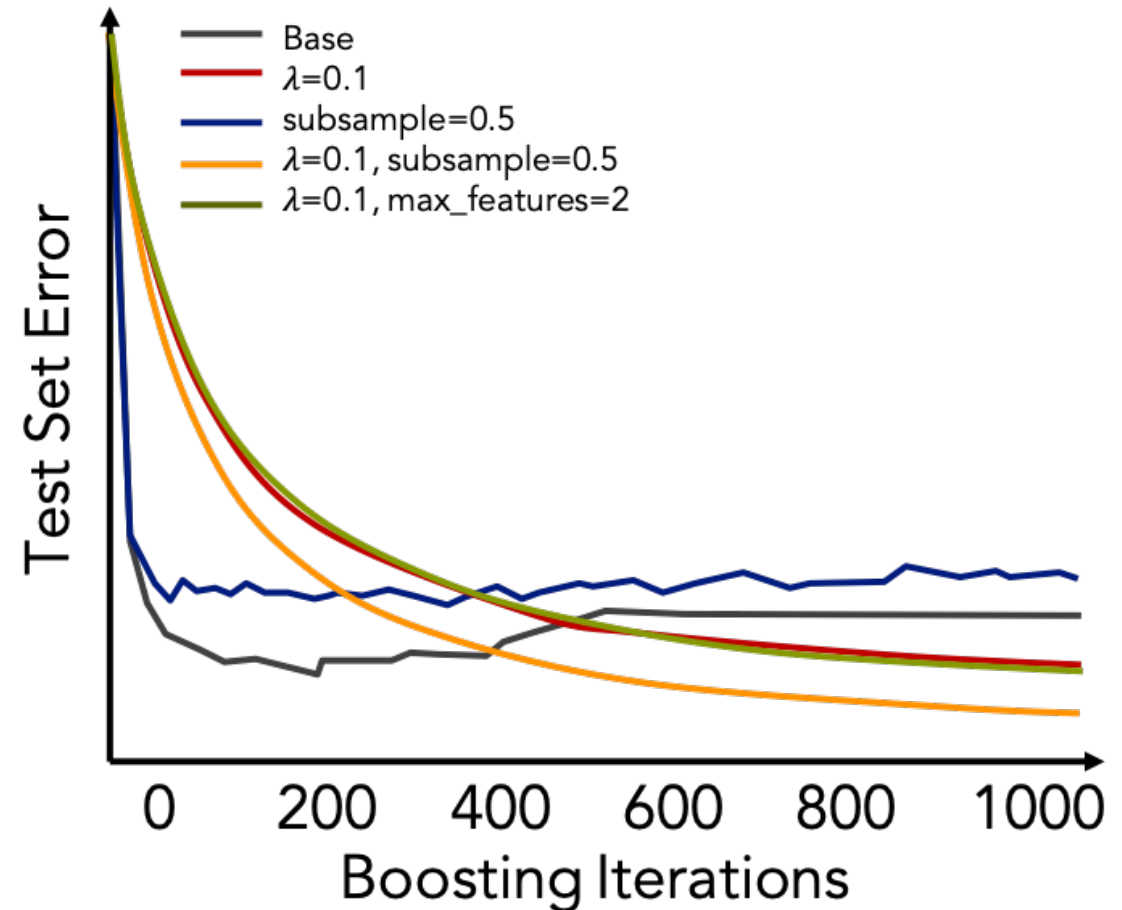
- Typically values are  $< 0.1$ , use 0.01 or 0.001
- Can choose  $M$  for early stopping





# STOCHASTIC GRADIENT BOOSTING

- Subsampling: each iteration, a fraction of training observations (without replacements) is used to grow the tree
- Can subsample on record or features
- Not only reduces computing time, can also produce more accurate model (less overfitting)



# BAGGING VS BOOSTING

## Bagging

- Bootstrapped samples
- Base trees created independently
- Only data points considered
- No weighting used
- Excess trees will not overfit

## Boosting

- Fit entire data set (can have subsamples)
- Base trees created successively
- Use residuals from previous models
- Up-weight misclassified points (explicit weight or as residuals)
- Beware of overfitting

# SKLEARN: GRADIENT BOOSTING TREE

```
>>> from sklearn.datasets import make_hastie_10_2
>>> from sklearn.ensemble import GradientBoostingClassifier
```

```
>>> X, y = make_hastie_10_2(random_state=0)
>>> X_train, X_test = X[:2000], X[2000:]
>>> y_train, y_test = y[:2000], y[2000:]
```

```
>>> clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,
...     max_depth=1, random_state=0).fit(X_train, y_train)
>>> clf.score(X_test, y_test)
0.913...
```

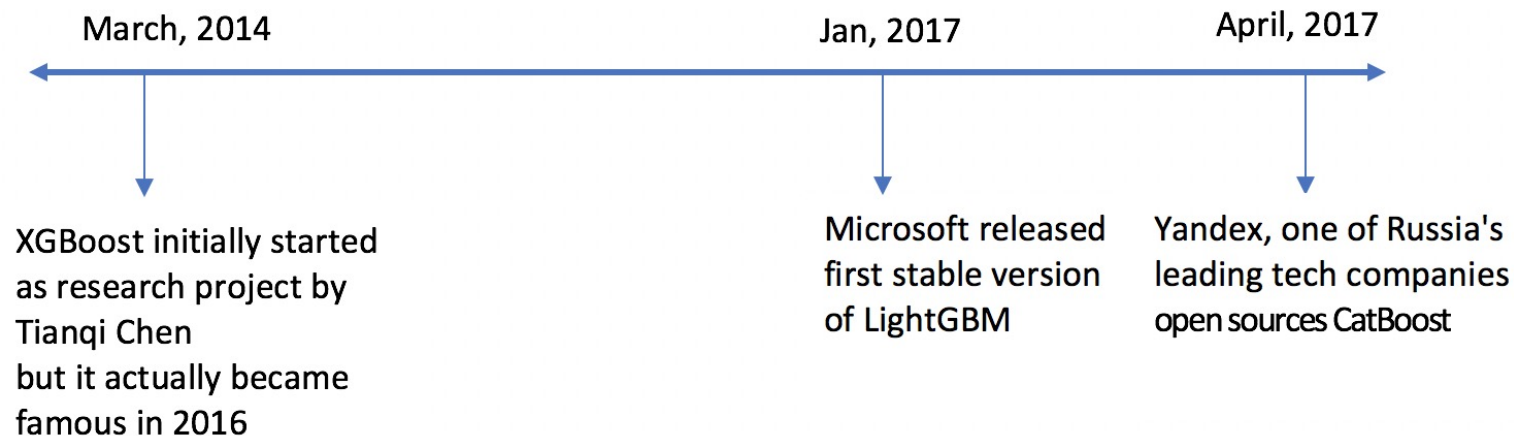
# GRADIENT BOOSTING LIBRARIES

- XGBoost (Extreme Gradient Boosting)

<https://xgboost.readthedocs.io>

- LightGBM (Light Gradient Boosted Machine)

<https://lightgbm.readthedocs.io/>



# XGBOOST AND LIGHTGBM

- Implements basic idea of GBM with some tweaks
  - Regularization of base tree
  - Approximate split finding
  - Weighted quantile sketch
  - Sparsity-aware split finding
  - Cache-aware block structure for out of core computation

# EXTREME GRADIENT BOOSTING (XGBOOST)

**Homesite Quote Conversion, Winners' Interview: 3rd place, Team New Model Army | CAD & QuY**

Kaggle Team | 02.29.2016



**Prudential Life Insurance Assessment, Winner's Interview: 2nd place, Bogdan Zhurakovskiy**

Kaggle Team | 03.14.2016



**Airbnb New User Bookings, Winner's Interview: 2nd place, Keiichi Kuroyanagi (@Keiku)**

Kaggle Team | 03.17.2016



**Telstra Network Disruption, Winner's Interview: 1st place, Mario Filho**

Kaggle Team | 03.23.2016



What these various data mining competitors have in common: all used XGBoost

# ENSEMBLE METHODS

- Ensemble of same classifiers
  - Bagging and random forest
  - Boosting and gradient boosted tree
- Ensemble of different classifiers