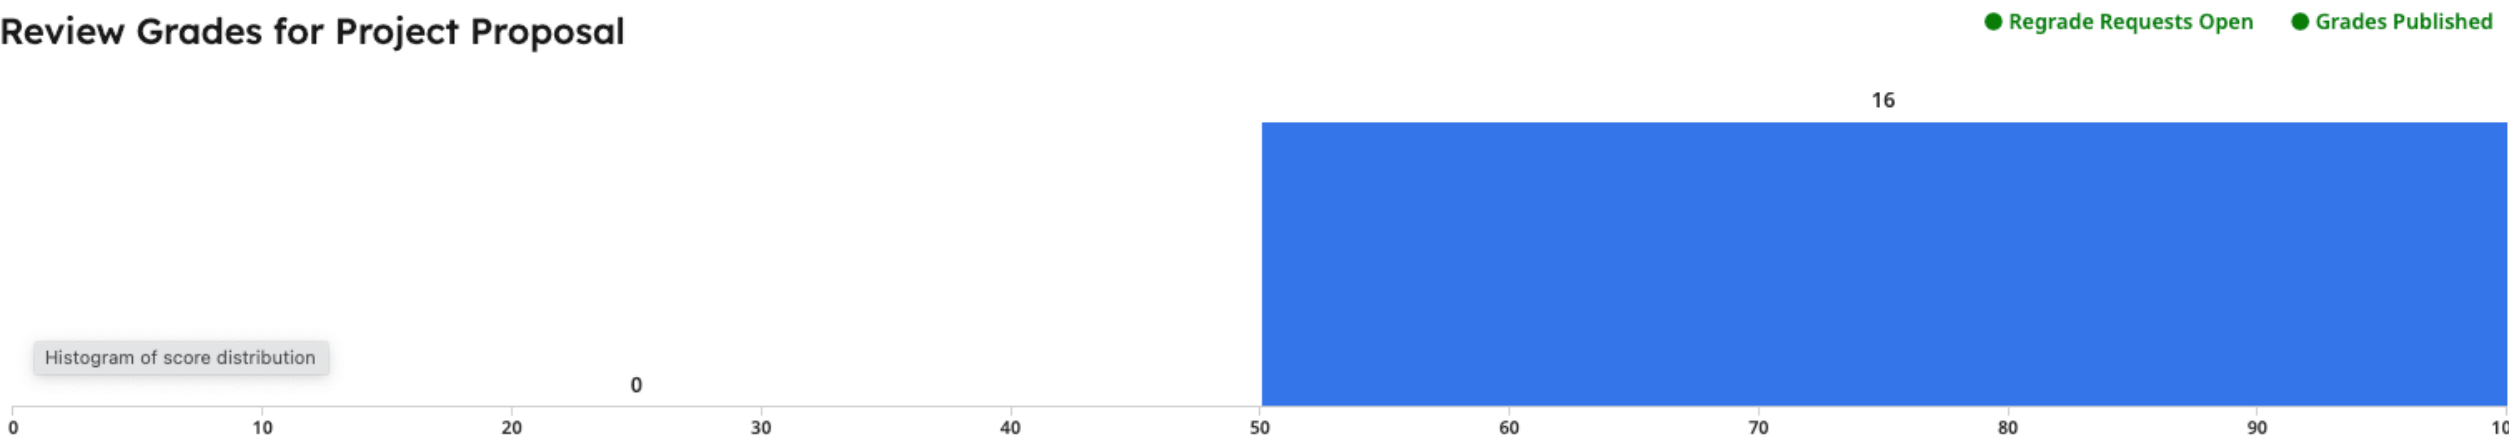


PROPOSAL GRADES POSTED

Review Grades for Project Proposal



Minimum

95.0

Median

98.5

Maximum

100.0

Mean

98.0

Std Dev [?](#)

1.51

SPOTLIGHT REMINDER

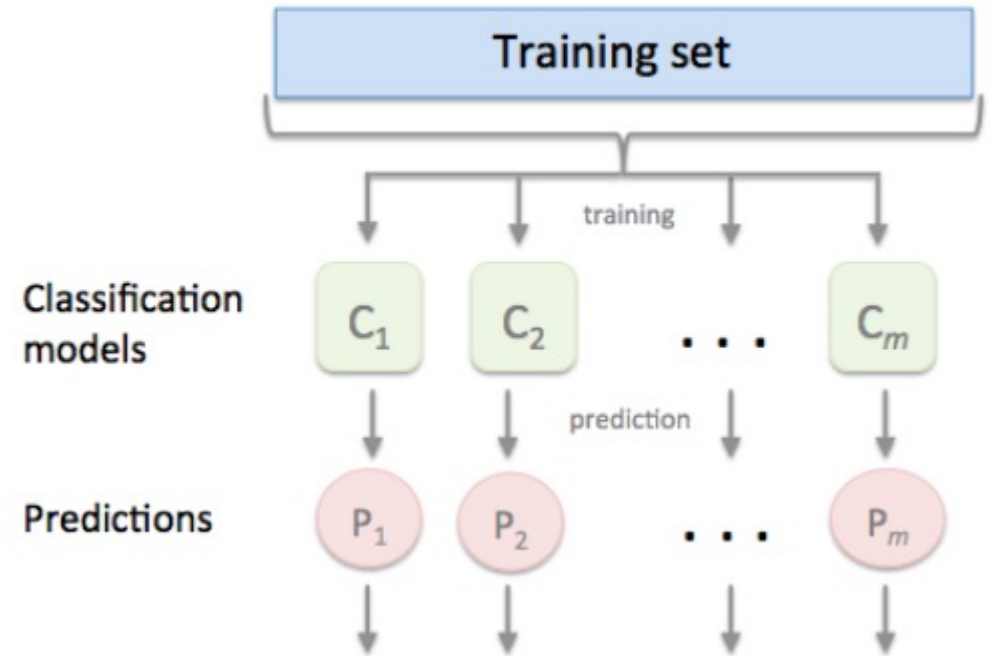
- Spotlight PPT slide due 10/30 on Canvas
- Project spotlight presentation on 11/1
 - 2 min each group
 - 1 min between groups
 - Group order 16 -> 1

ENSEMBLE METHODS

- Ensemble of same classifiers
 - Bagging and random forest
 - Boosting and gradient boosted tree
- Ensemble of different classifiers

VOTING

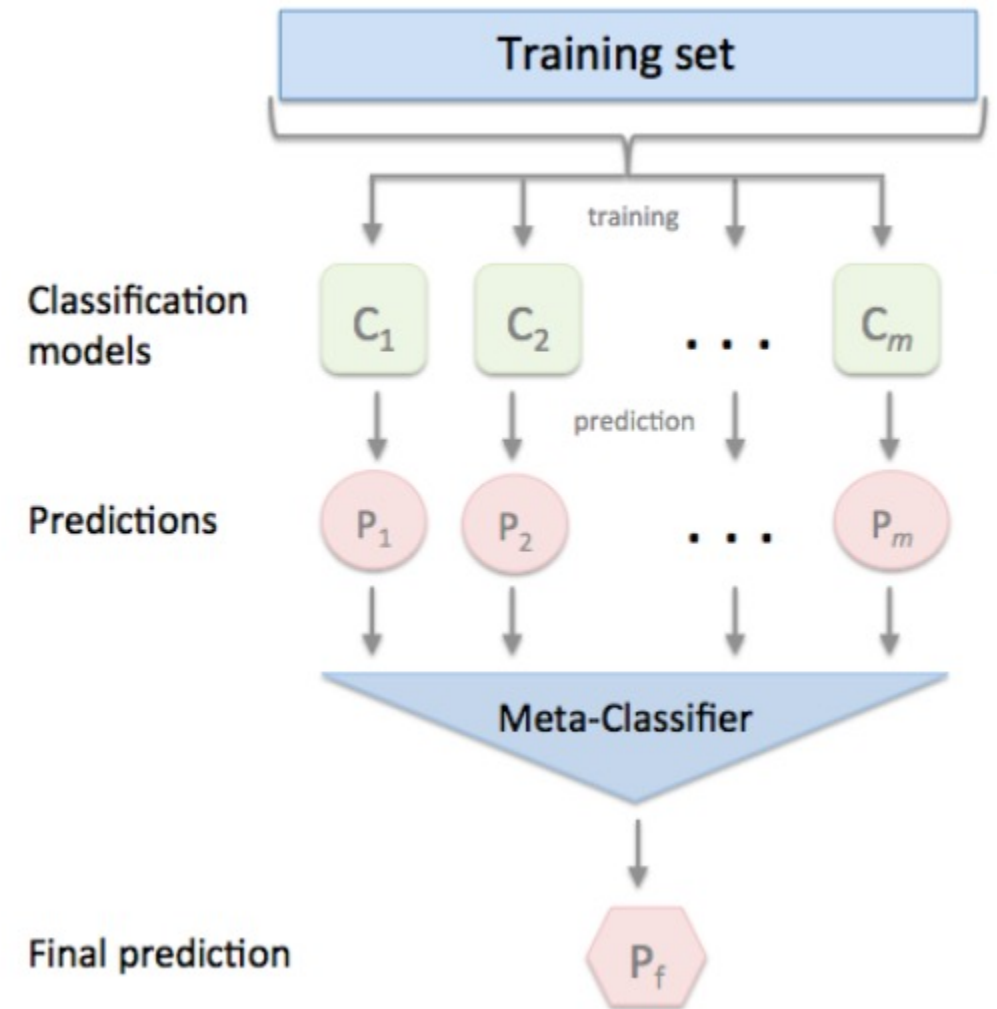
- Voting ensembles mimic error-correcting codes
 - More voters \longrightarrow potential better signal to noise
 - Lower correlation between models



How do we combine the predictions?

STACKING

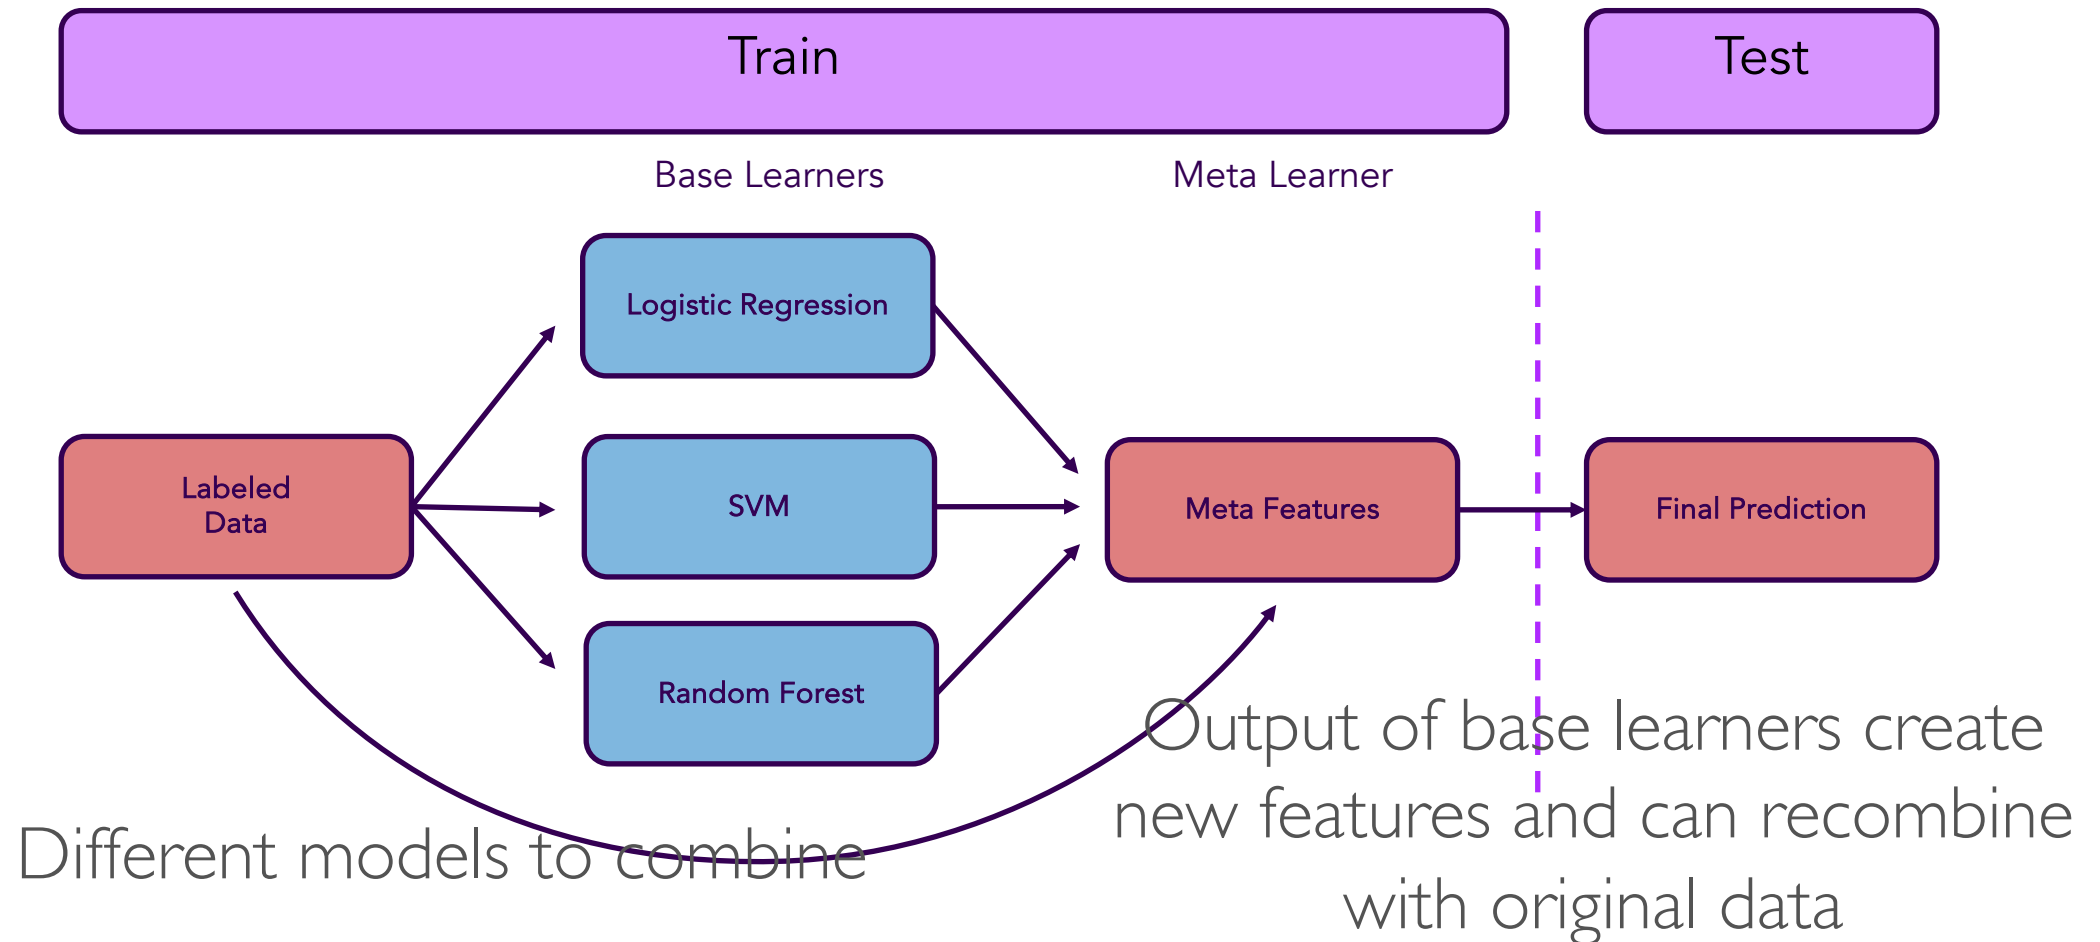
- Introduced by Wolpert, 1992
- Use a pool of base classifiers and do **out-of-fold** predictions, then train a meta-classifier to combine their predictions
- Stacker model gains information by using first-stage predictions as features



BLENDING

- Close to stacked generalization, but a bit simpler
- Instead of out-of-fold predictions, create small **holdout set** that the stacker is then trained on this set
- Disadvantages:
 - Less data used overall
 - Final model may overfit to holdout set

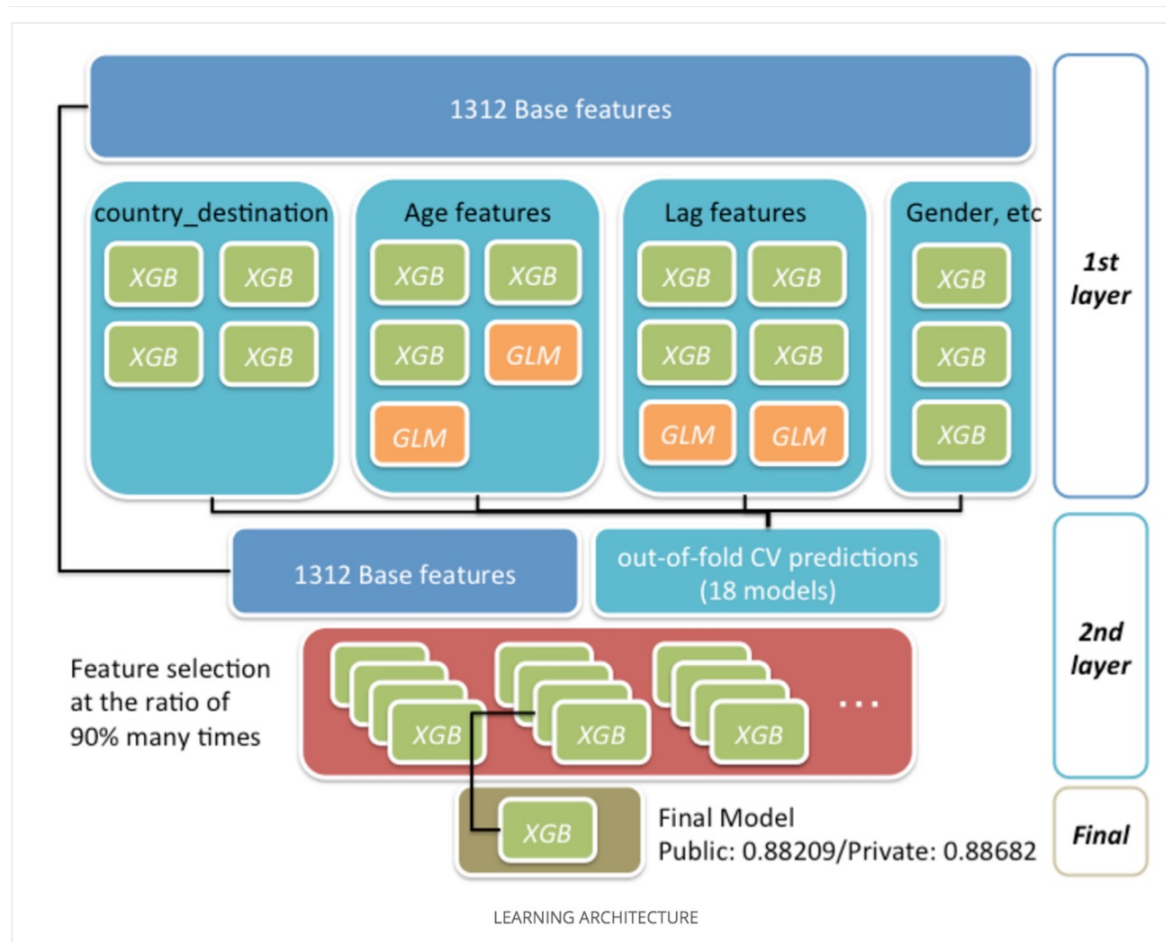
EXAMPLE: COMBINING DIFFERENT CLASSIFIERS



STACKING AND BLENDING

- Why stop at two stages? Why not combine multiple ensemble models?

EXAMPLE: AIRBNB 2ND PLACE

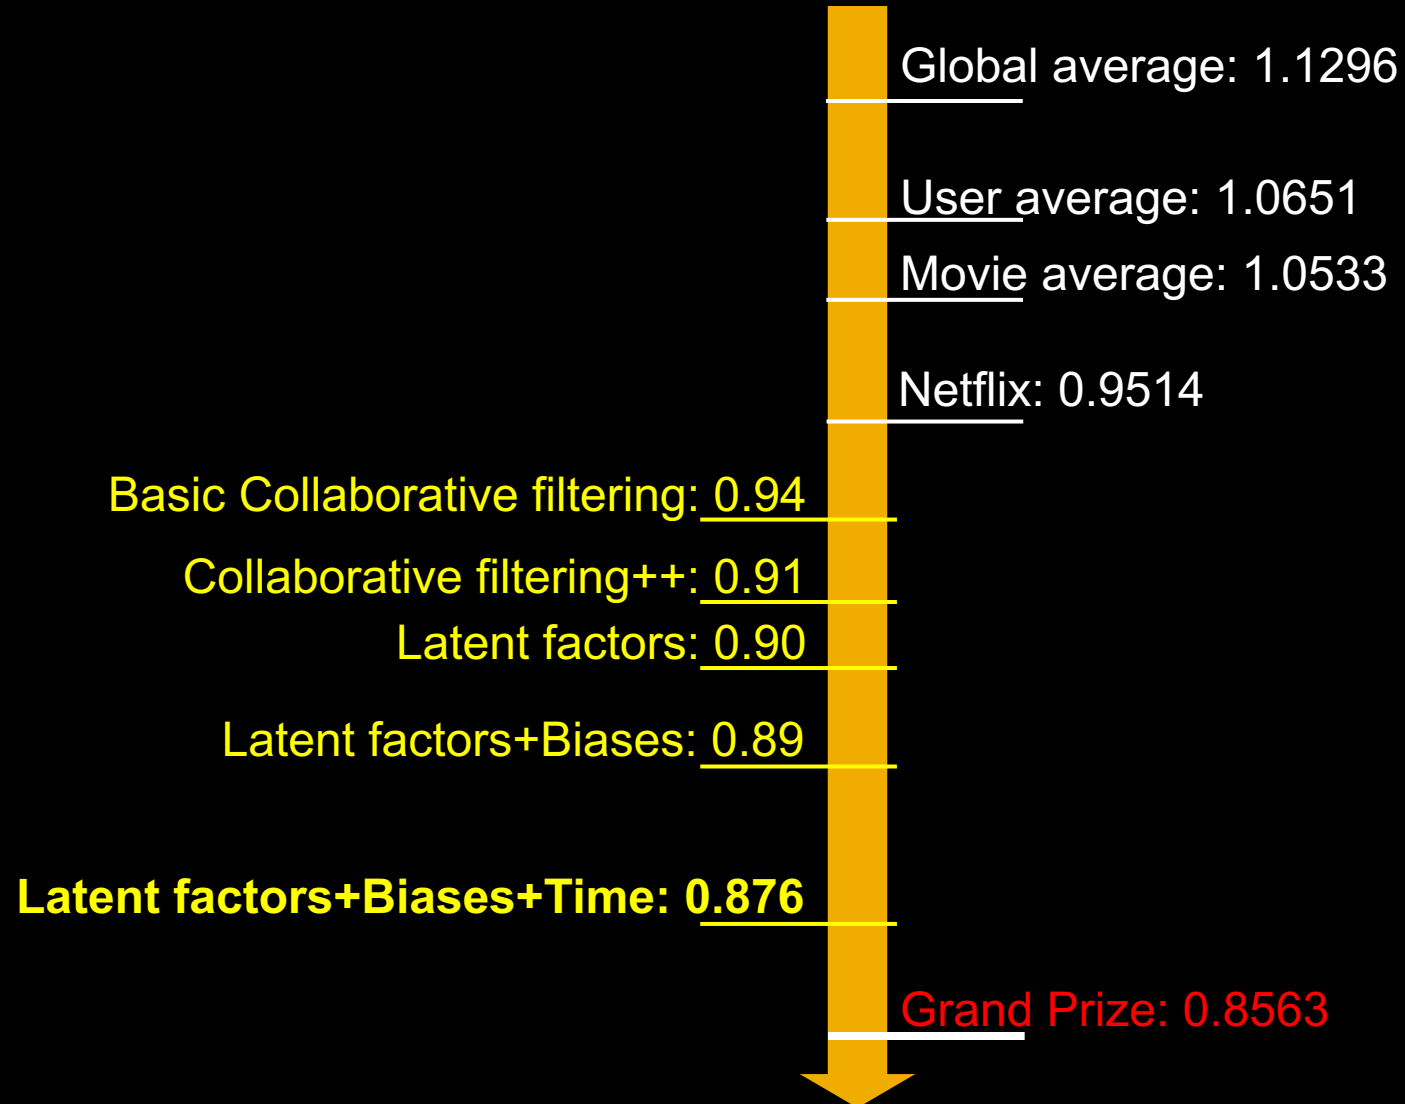


<https://github.com/Keiku/kaggle-airbnb-recruiting-new-user-bookings>

The Netflix Prize

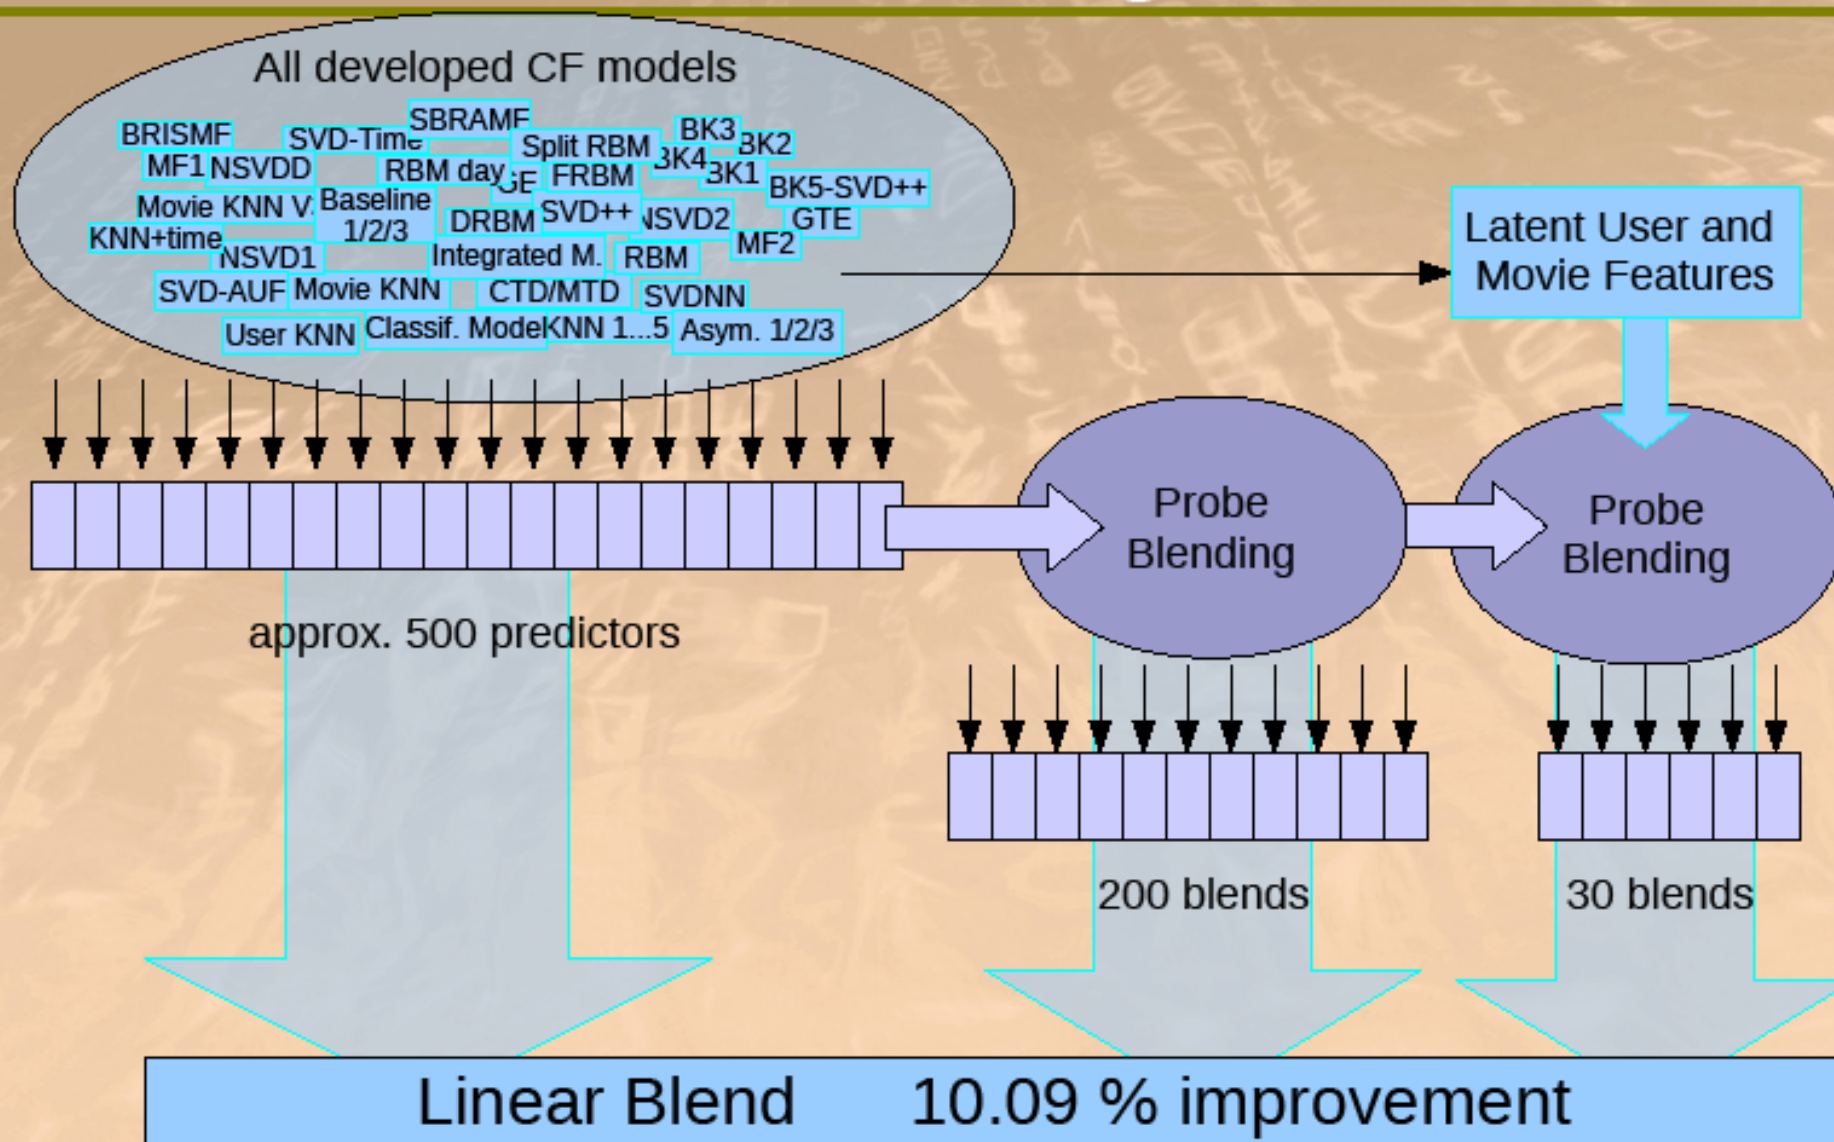
- An open competition 2006 – 2009 held by Netflix
- **Training data**
 - 100 million ratings (with timestamps), 480,000 users, 17,770 movies
 - 6 years of data: 2000-2005
- **Test data**
 - Last few ratings of each user (2.8 million)
- **Evaluation criterion: Root Mean Square Error**
$$(\text{RMSE}) = \frac{1}{|R|} \sqrt{\sum_{(i,x) \in R} (\hat{r}_{xi} - r_{xi})^2}$$
- **Competition**
 - Netflix's system Cinematch RMSE: **0.9514**
 - **\$1 million** prize for 10% improvement on Netflix
 - \$50,000 progress prize every year

Netflix Prize



The big picture

Solution of BellKor's Pragmatic Chaos



DIMENSIONALITY REDUCTION

CS 334: Machine Learning

TYPES OF UNSUPERVISED LEARNING

Clustering

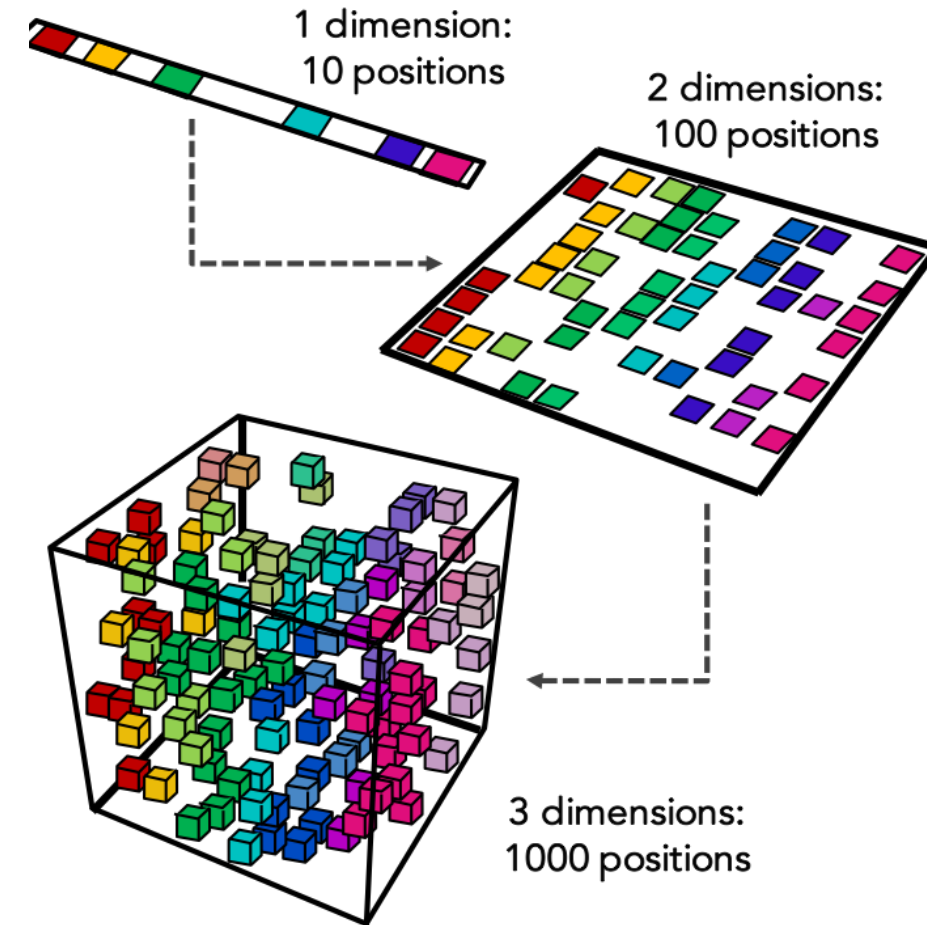
identify unknown structure in the data

Dimensionality Reduction

use structural characteristics to simplify data

CURSE OF DIMENSIONALITY

- Increasing features should improve performance right?
- In practice, too many features leads to worse performance
- # of training examples required increases exponentially with the # of features

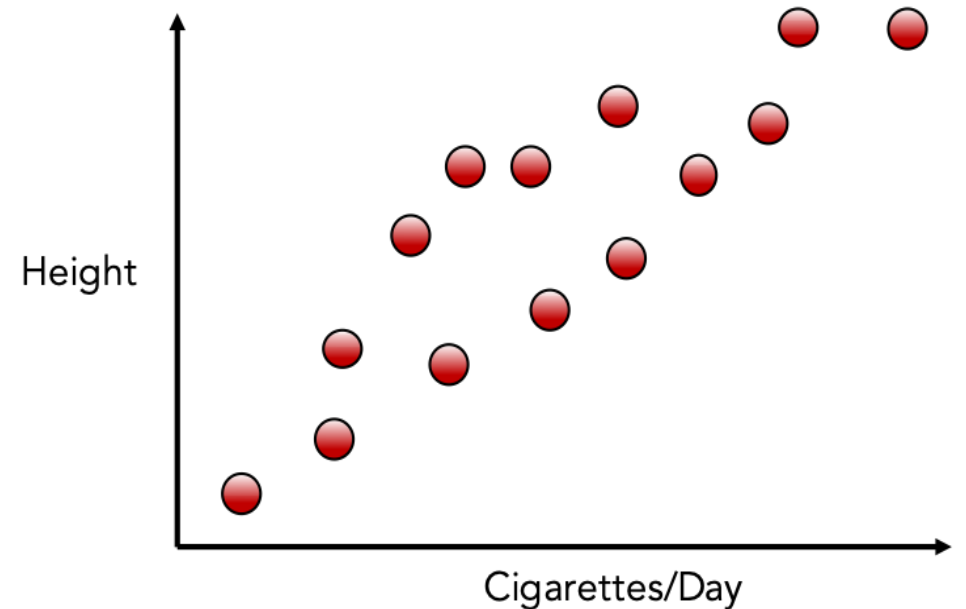


FEATURE SELECTION/REDUCTION

- Feature selection
 - Filter methods (agnostic to the learning algorithm)
 - Wrapper methods (keep model in the loop)
 - Embedded methods (e.g. Lasso regularization)
- Dimension reduction

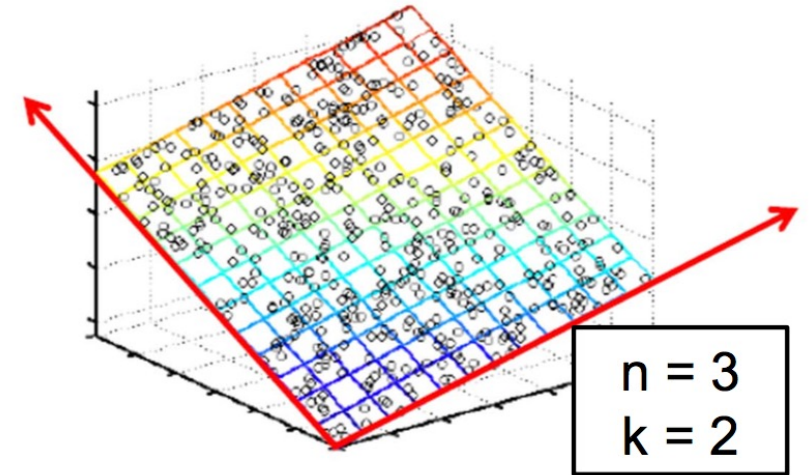
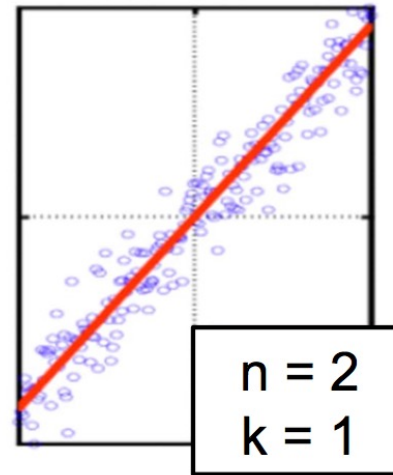
MOTIVATION: EXAMPLE

- Two features: Height and cigarettes per day
- Both features are correlated
- Can we reduce the features to one?



DIMENSIONALITY REDUCTION

- Represent data with fewer dimensions
- Discover “intrinsic dimensionality” of data
- New feature space: linear / non-linear combination of original features



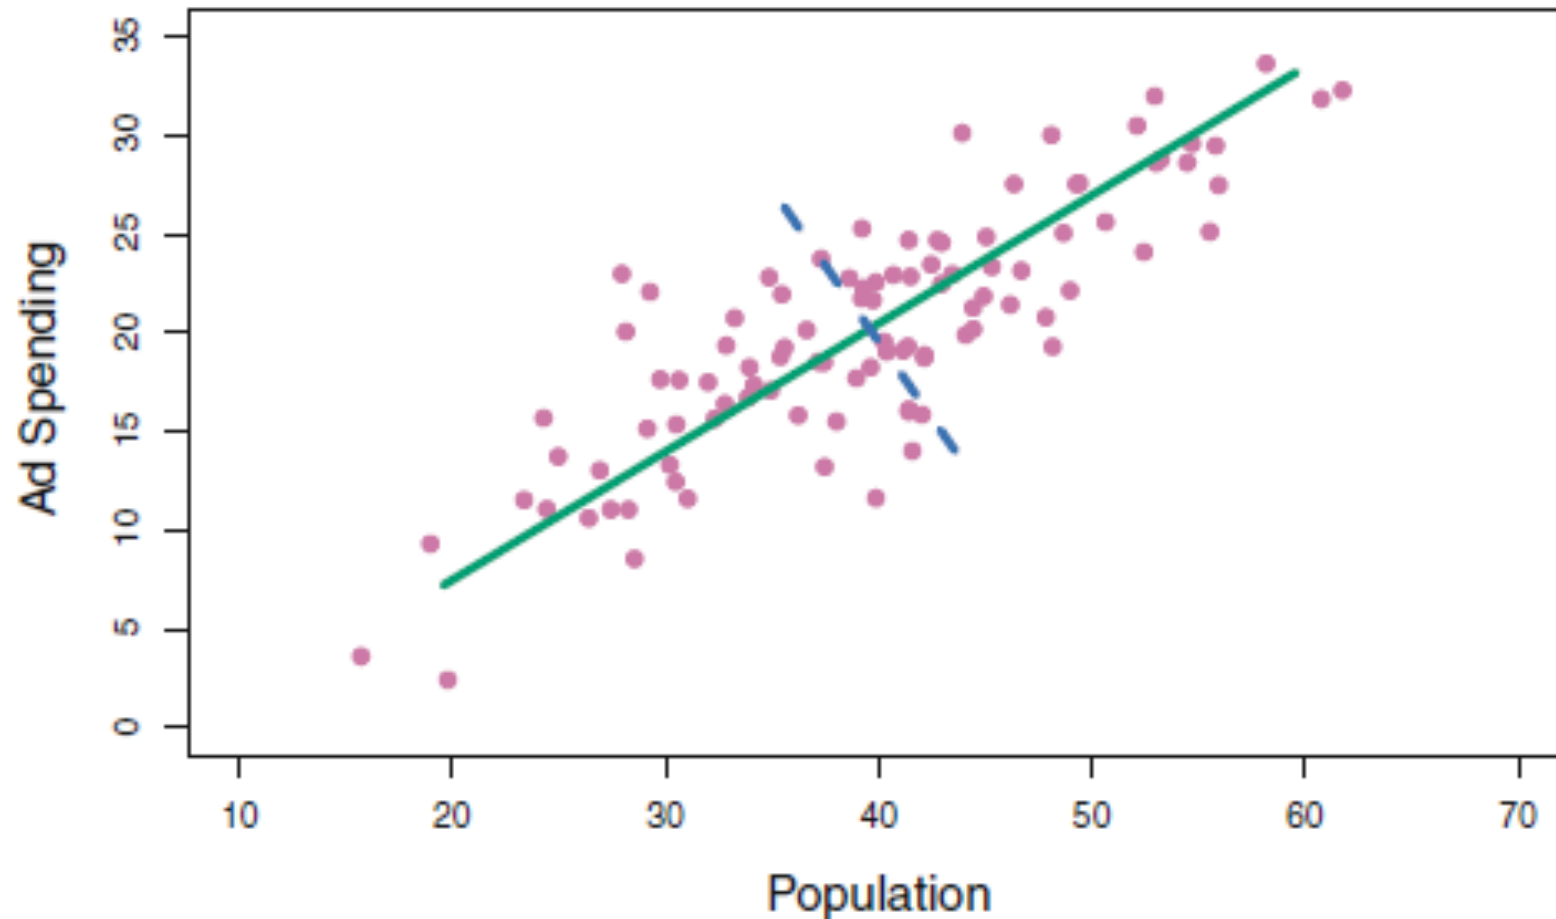
WHY DIMENSIONALITY REDUCTION

- Simplified data processing
- Noise reduction (removing feature redundancy)
- Easier learning — less parameters
- Robust learning — numerical stability due to less correlations
- Easier visualization — show high dimensional data in 2D or 3D



GROUP ACTIVITY

WHICH IS THE BEST PROJECTION AND WHY?

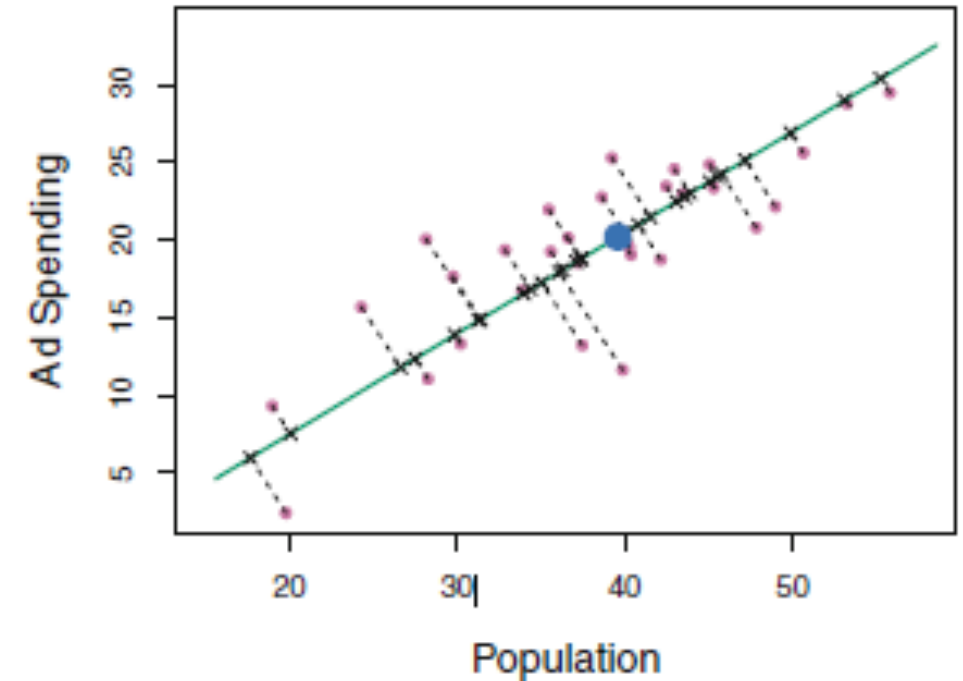


PRINCIPAL COMPONENT ANALYSIS (PCA)

- Developed by Pearson in 1901
- Popular and widely studied
- Finds sequence of linear combinations of the features (also known as principal components) that have maximal variance and are uncorrelated

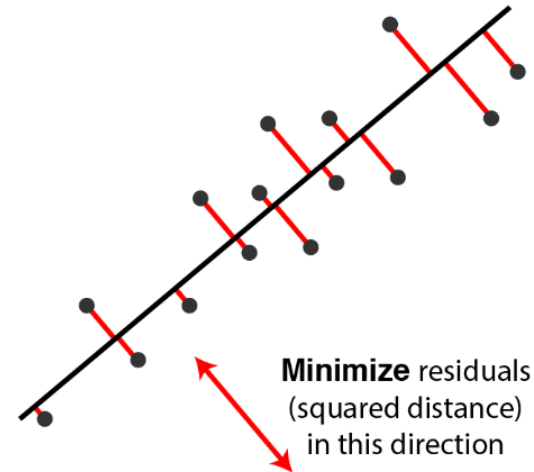
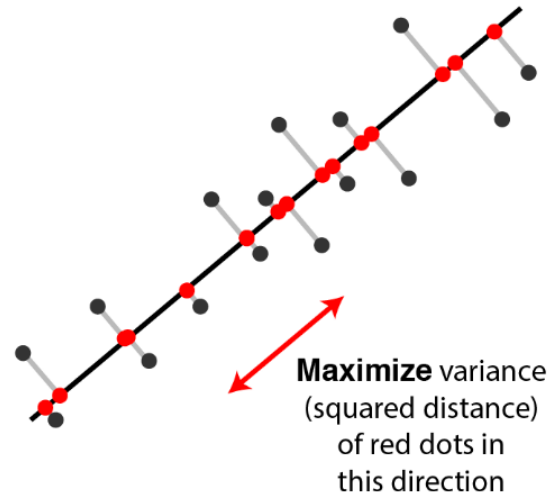
PRINCIPAL COMPONENT ANALYSIS (PCA)

- First principal component
 - Yields the highest variance of the projection
 - Minimizes sum of squared perpendicular distances (reconstruction error between data and projected points)

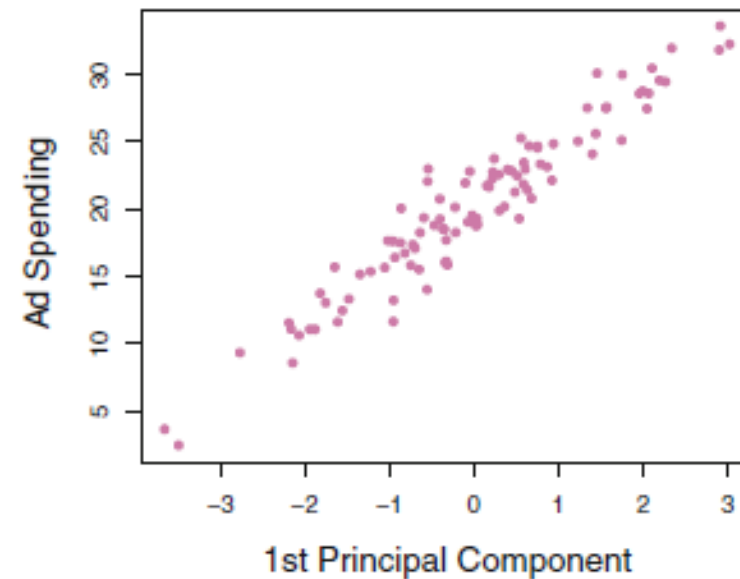
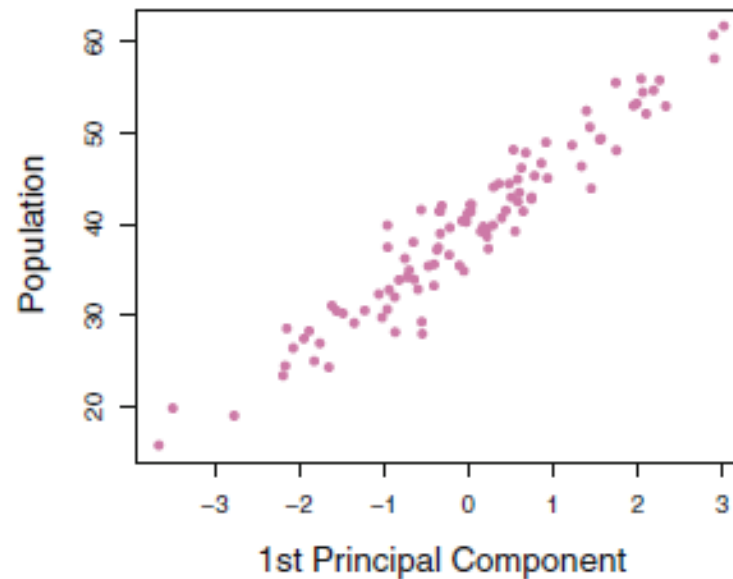
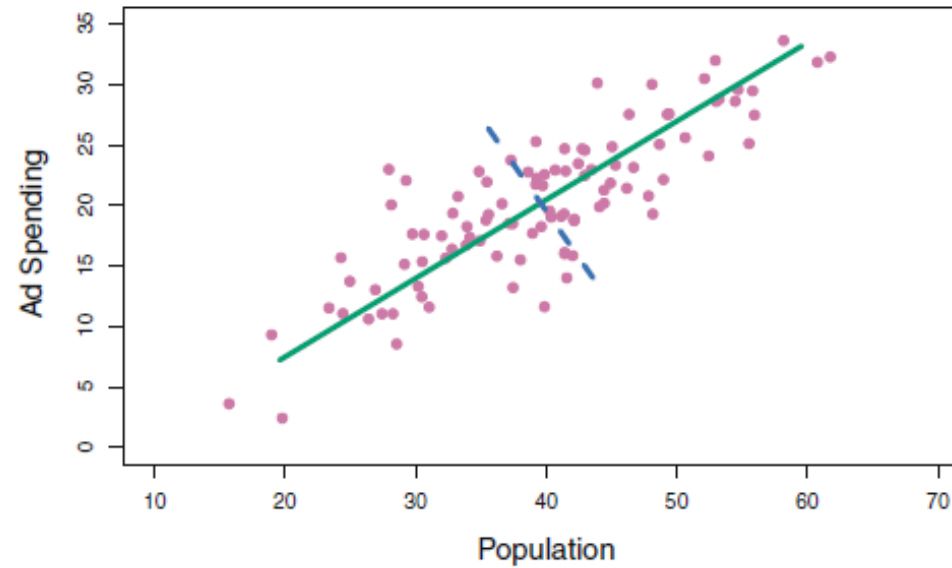


TWO INTERPRETATIONS OF PCA

- Maximize the variance of projection along principal component
- Minimize the reconstruction error between original data and projected components

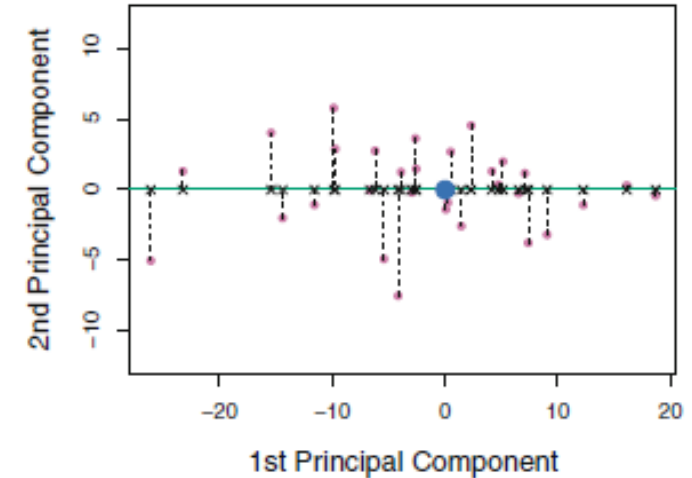
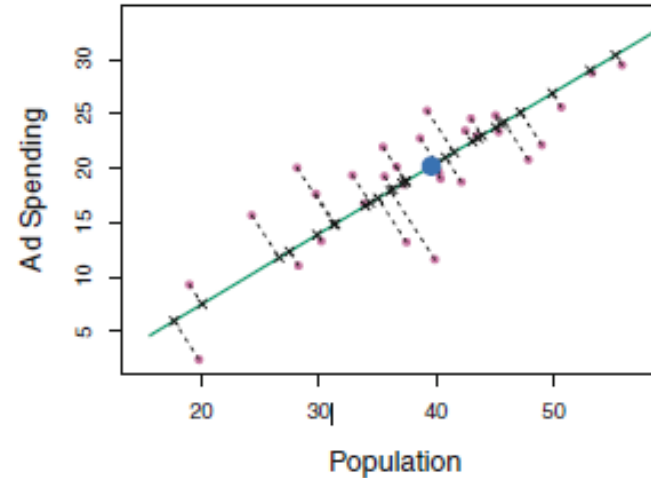


1ST PRINCIPAL COMPONENT

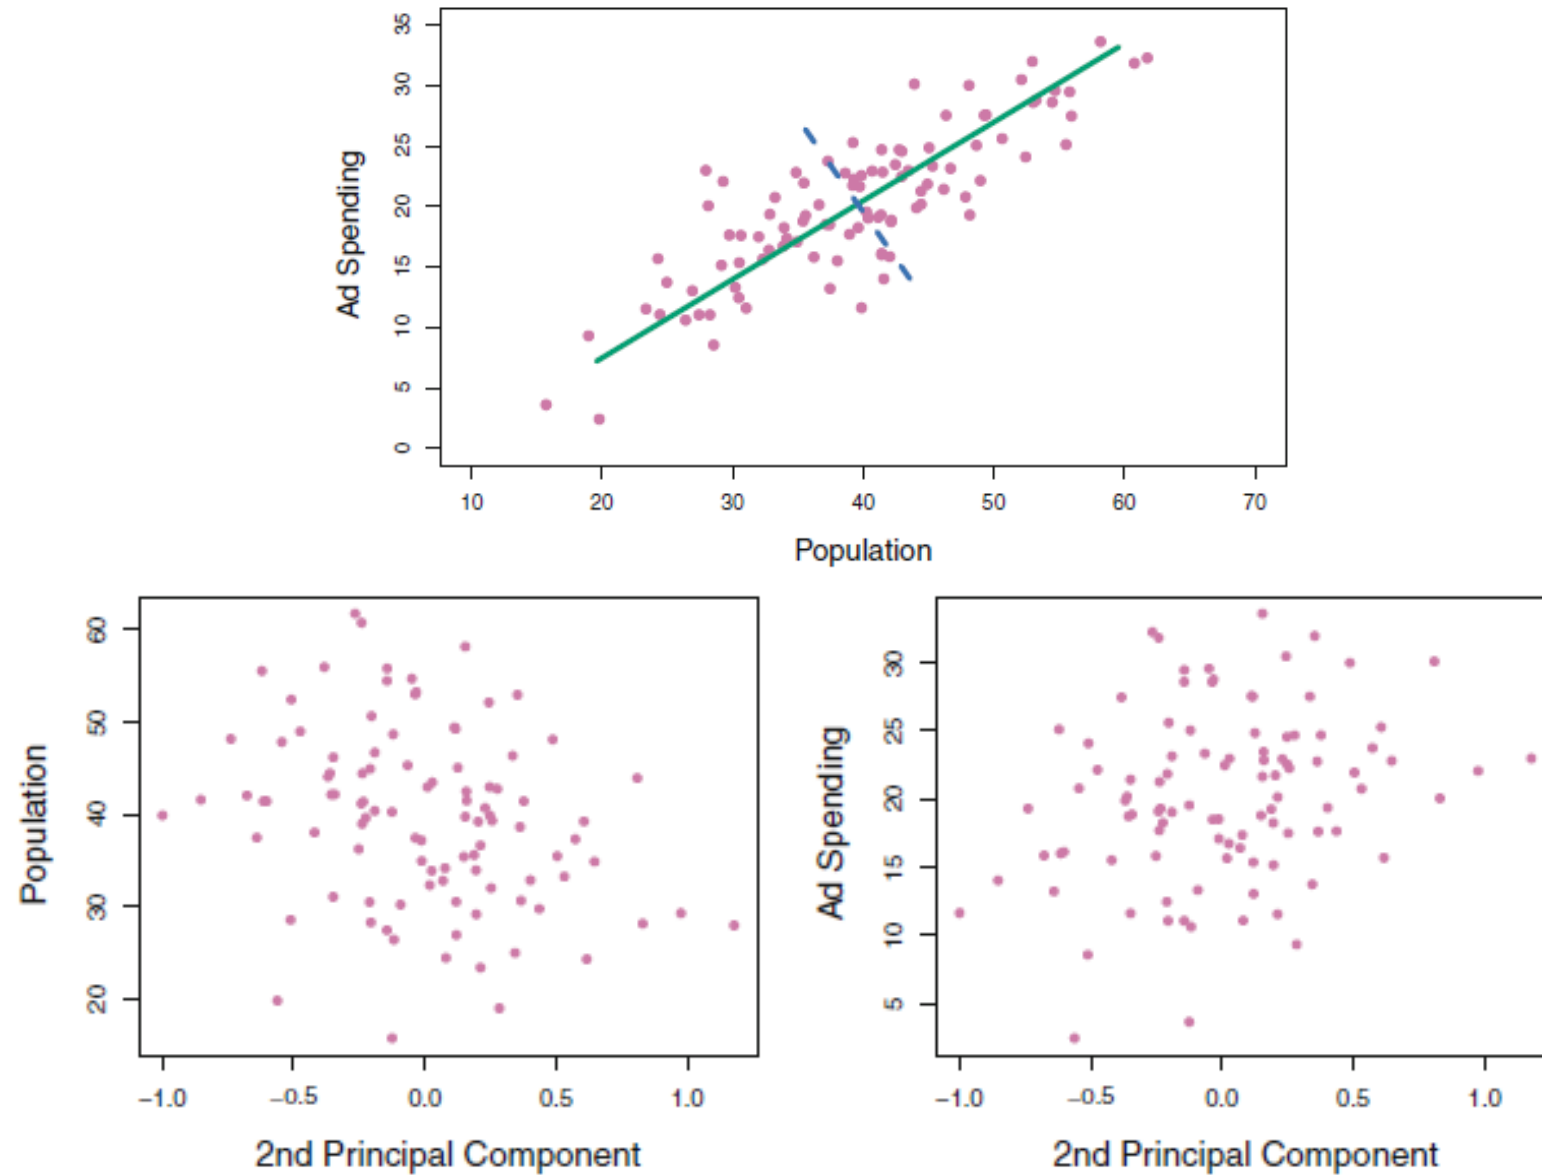


2ND PRINCIPAL COMPONENT

- Second principal component
 - Orthogonal to the first principal component
 - And has largest variance
- In general, we can construct up to p principal components (p features)



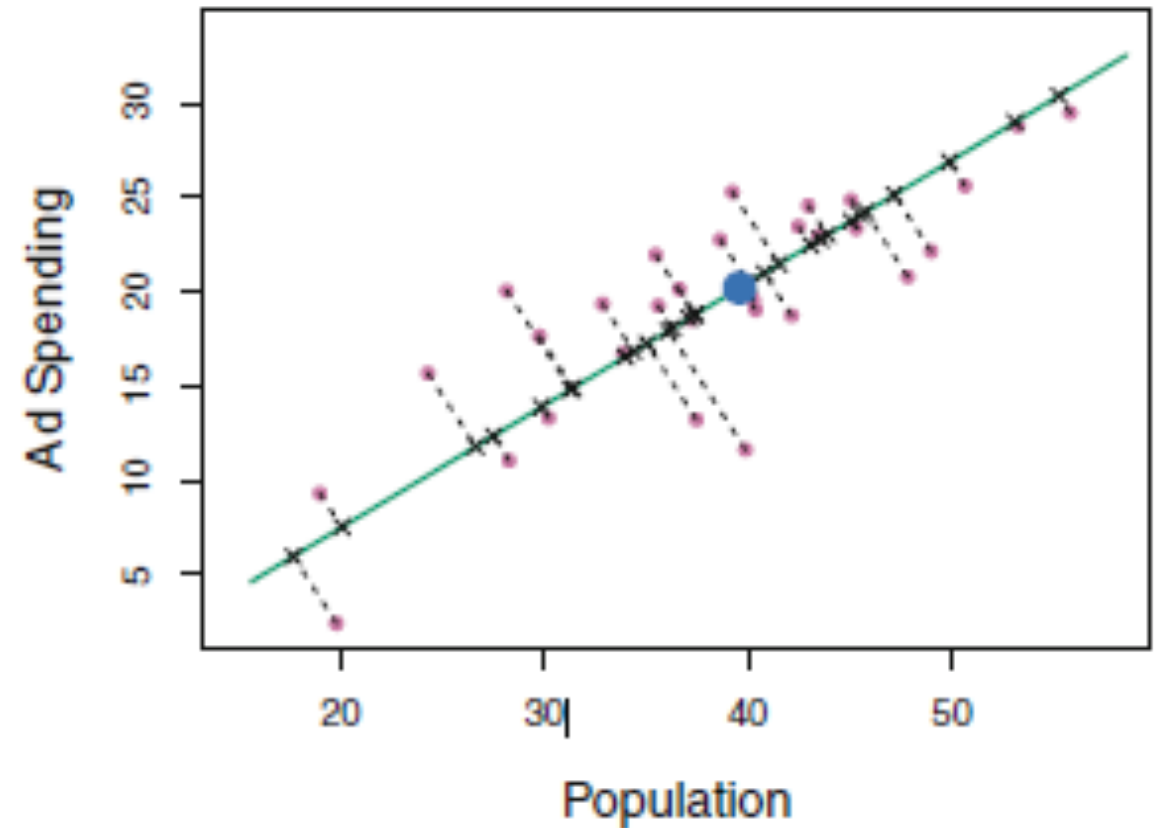
2ND PRINCIPAL COMPONENT



PRINCIPAL COMPONENT ANALYSIS (PCA)

- Principal component loadings
 - Determines the axis
 - $v = (0.839, 0.544)$
- Principal component scores
 - Projected position of each point on the axis
 - $Z = Xv$

(Data is centered first)



$$Z_1 = 0.839 \times (\text{pop} - \overline{\text{pop}}) + 0.544 \times (\text{ad} - \overline{\text{ad}}).$$

PROJECTION ONTO UNIT VECTORS

- Definition of dot product:

$$\mathbf{A} \cdot \mathbf{B} = \|\mathbf{A}\|_2 \|\mathbf{B}\|_2 \cos \theta$$

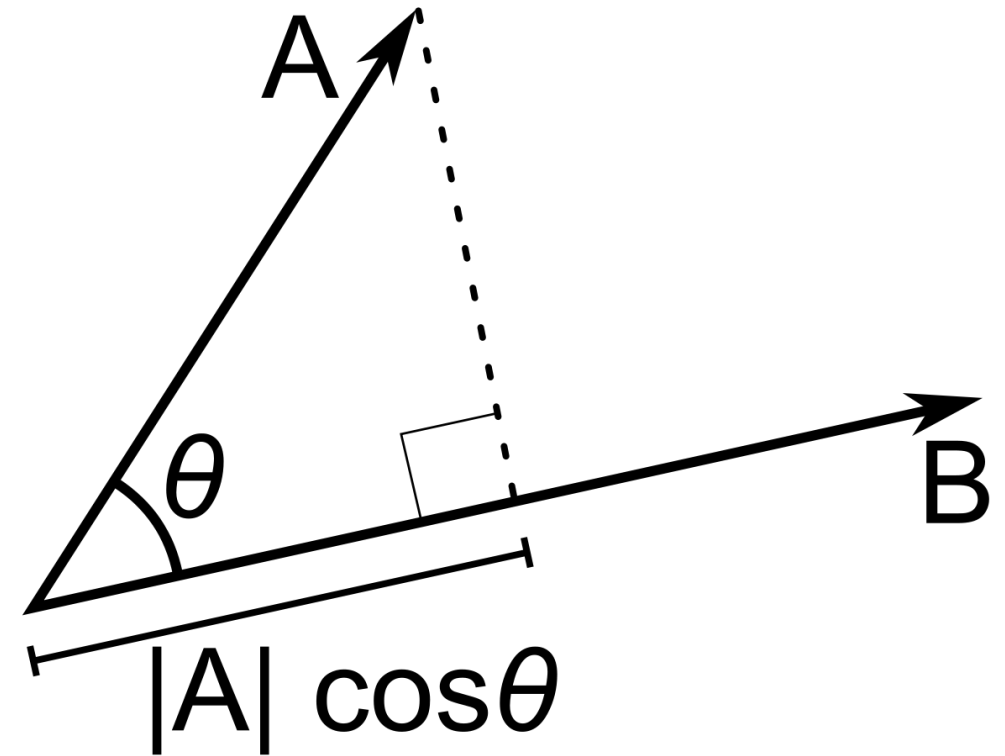
- If B is a unit vector, dot product is length of the projection

$$\mathbf{A} \cdot \mathbf{B} = \|\mathbf{A}\|_2 \cos \theta$$

- Projection of A onto B (vector):

$$(\mathbf{A} \cdot \mathbf{B}) \mathbf{B}$$

Coefficient / score



https://en.wikipedia.org/wiki/Dot_product

EXAMPLE: US ARRESTS DATASET

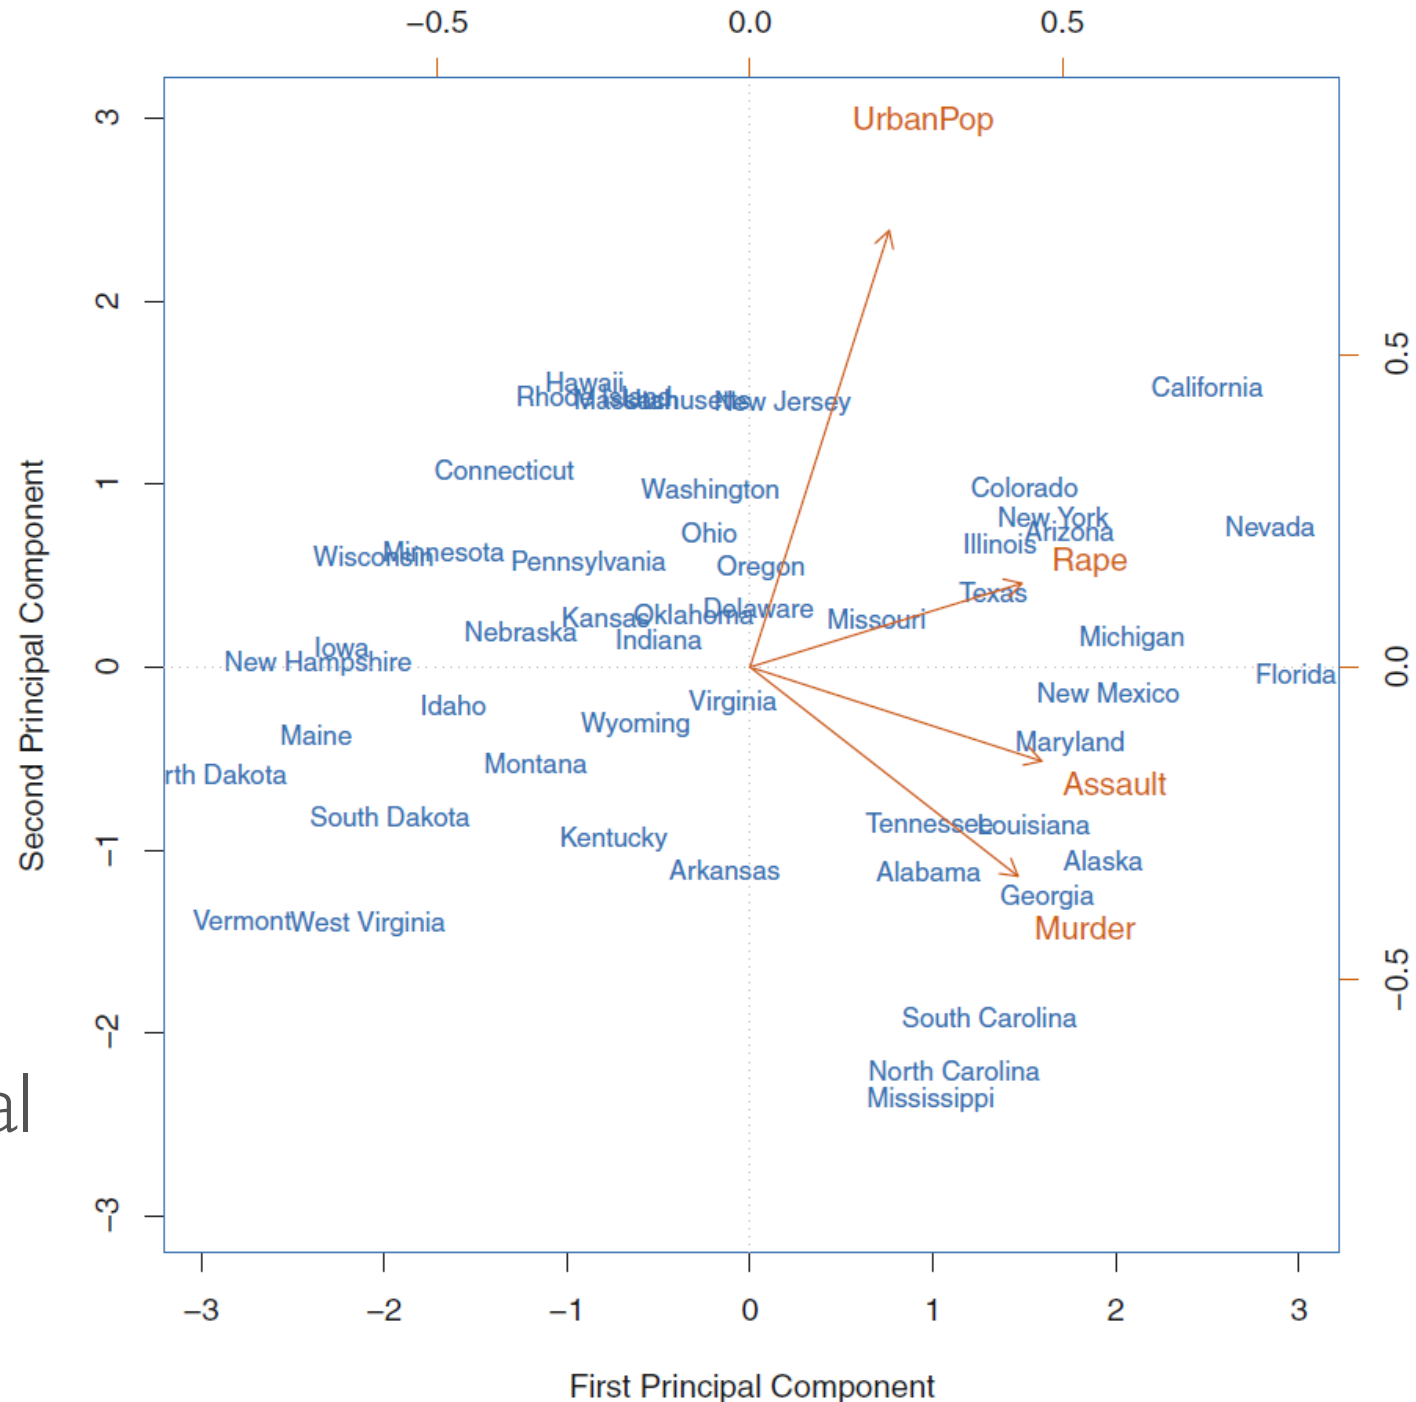
- 50 states with 4 features
 - Number of arrests per 100,000 residents for Murder, Assault, Rape
 - Percent of population living in urban areas (UrbanPop)

1		Murder	Assault	UrbanPop	Rape
2	Alabama	13.2	236	58	21.2
3	Alaska	10	263	48	44.5
4	Arizona	8.1	294	80	31
5	Arkansas	8.8	190	50	19.5
6	California	9	276	91	40.6
7	Colorado	7.9	204	78	38.7
8	Connecticut	3.3	110	77	11.1
9	Delaware	5.9	238	72	15.8
10	Florida	15.4	335	80	31.9
11	Georgia	17.4	211	60	25.8

EXAMPLE: US ARRESTS (BIPLOT)

	PC1	PC2
Murder	0.5358995	-0.4181809
Assault	0.5831836	-0.1879856
UrbanPop	0.2781909	0.8728062
Rape	0.5434321	0.1673186

Biplot: displays both principal component scores and principal component loadings.



PCA: 1st PC

- 1st PC of X is unit vector that maximizes the sample variance compared to all other unit vectors

$$\mathbf{v}_1 = \operatorname{argmax}_{\|\mathbf{v}\|_2=1} (\mathbf{X}\mathbf{v})^\top (\mathbf{X}\mathbf{v})$$

- 1st PC score: $\mathbf{X}\mathbf{v}_1$
- Variance explained by first PC: $(\mathbf{X}\mathbf{v}_1)^\top (\mathbf{X}\mathbf{v}_1)/n$

PCA: NEXT PC

- Idea: Successively find orthogonal directions of highest variance
- Why orthogonal?
 - Want to minimize redundancy
 - Want to look at variance in different direction
 - Computation is easier

PCA: 2ND PC

- 2nd PC of \mathbf{X} is unit vector that is orthogonal to the 1st PC such that it maximizes the sample variance compared to all other unit vectors that are orthogonal to the 1st PC

$$\mathbf{v}_2 = \operatorname{argmax}_{\|\mathbf{v}\|_2=1, \mathbf{v}^\top \mathbf{v}_1=0} (\mathbf{X}\mathbf{v})^\top (\mathbf{X}\mathbf{v})$$

- 2nd PC score: $\mathbf{X}\mathbf{v}_2$
- Variance explained by 2nd PC: $(\mathbf{X}\mathbf{v}_2)^\top (\mathbf{X}\mathbf{v}_2)/n$

PCA: PROBLEM FORMULATION

- Recall 1st and 2nd PC

$$\mathbf{v}_1 = \underset{\|\mathbf{v}\|_2=1}{\operatorname{argmax}} (\mathbf{X}\mathbf{v})^\top (\mathbf{X}\mathbf{v})$$

$$\mathbf{v}_2 = \underset{\|\mathbf{v}\|_2=1, \mathbf{v}^\top \mathbf{v}_1=0}{\operatorname{argmax}} (\mathbf{X}\mathbf{v})^\top (\mathbf{X}\mathbf{v})$$

- Find orthonormal vectors that maximizes variance (assuming \mathbf{X} is zero-centered)

PCA: PROBLEM FORMULATION

- Given a feature matrix \mathbf{X} with n data points, find \mathbf{W} such that $\|\mathbf{W}\|_2 = 1$ and the $\text{Var}(\mathbf{XW})$ is maximized and \mathbf{W} consists of orthonormal vectors

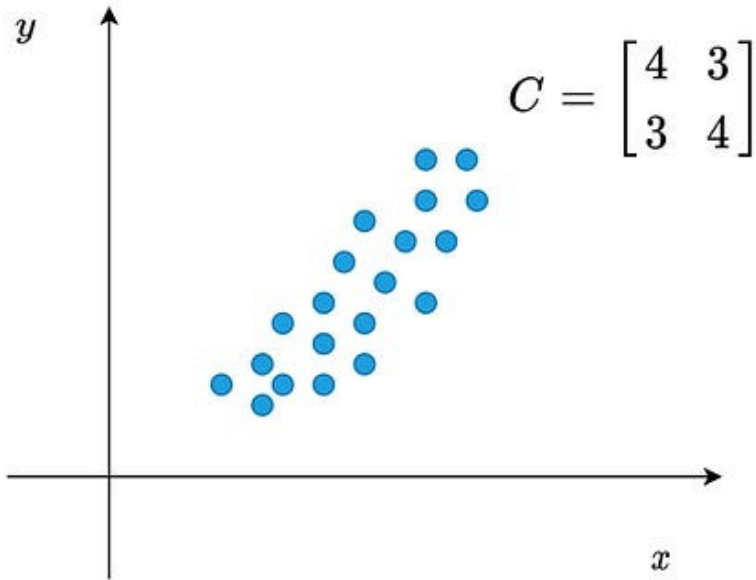
$$\begin{aligned}\text{Var}(\mathbf{XW}) &= \frac{1}{N} (\mathbf{W}^\top (\mathbf{X} - \mu_{\mathbf{X}})^\top (\mathbf{X} - \mu_{\mathbf{X}}) \mathbf{W}) \\ &= \mathbf{W}^\top \Sigma_{\mathbf{X}} \mathbf{W}\end{aligned}$$

Sample covariance matrix

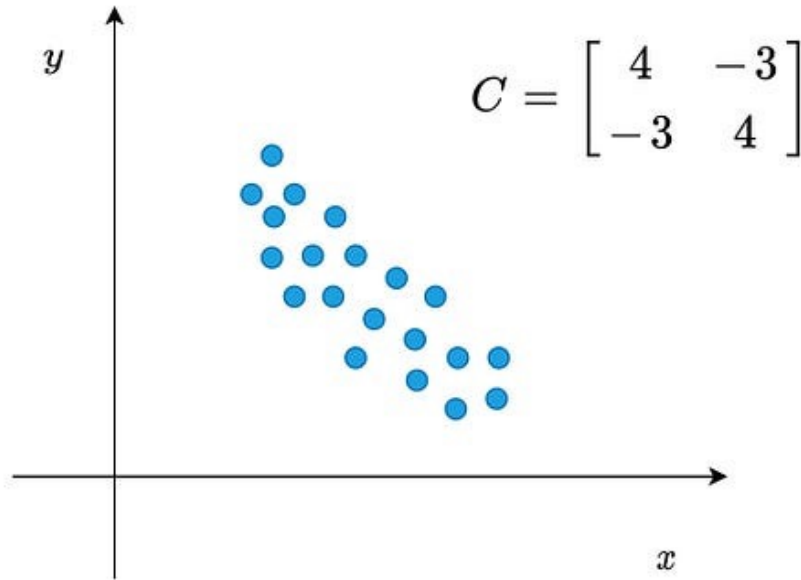
- The solution is the Eigenvectors of the covariance matrix

EXAMPLE: COVARIANCE MATRIX

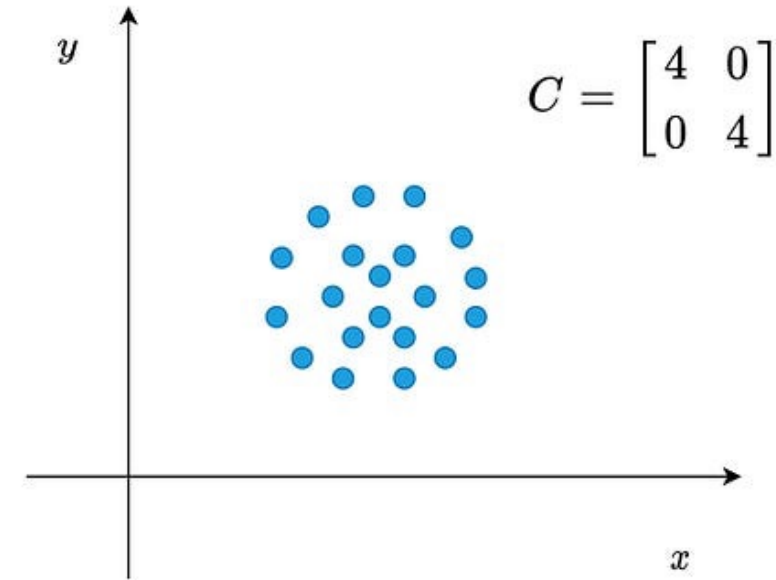
Positive
Covariance



Negative
Covariance

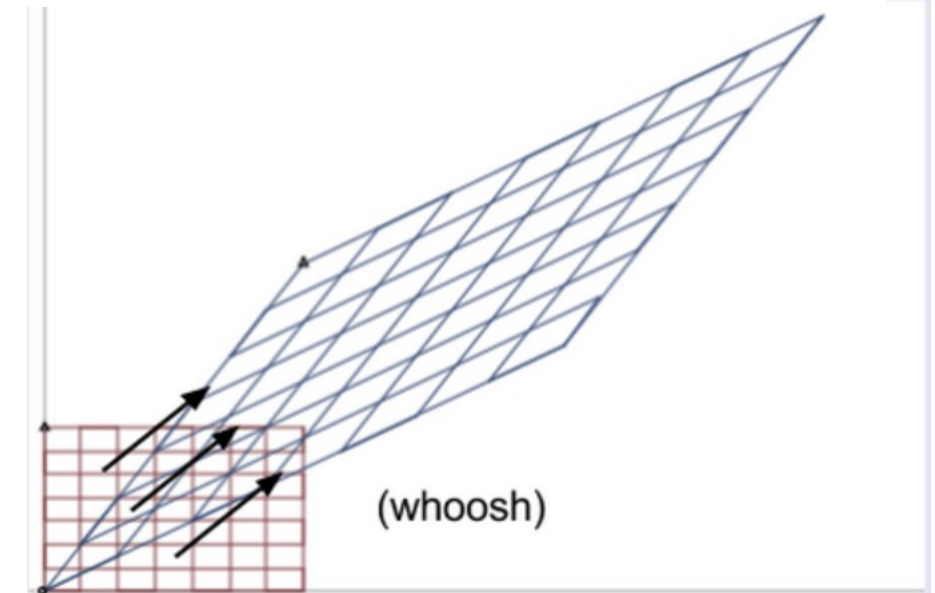
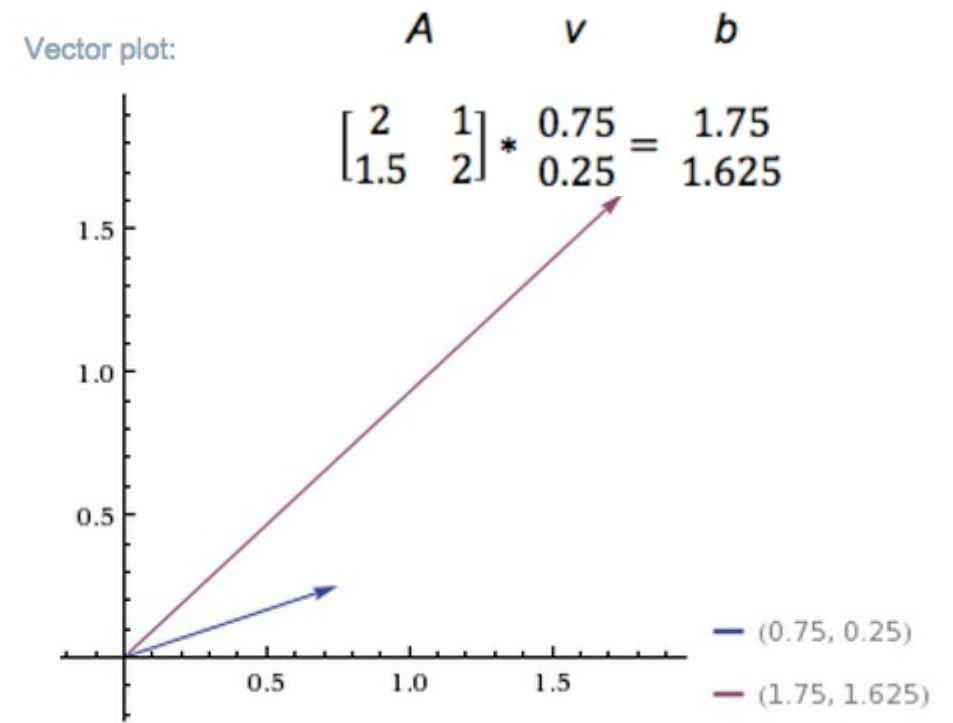


Zero
Covariance



REVIEW: EIGENVALUES + EIGENVECTORS

- Eigenvector and eigenvalue $\mathbf{Ax} = \lambda\mathbf{x}$
- Analogy: Matrix is a gust of wind (invisible force with visible result)
 - Eigenvector is like a weathervane which tells you the direction the wind is blowing in
 - Eigenvalue is just the scalar coefficient
- Eigenvectors of a symmetric matrix are orthonormal

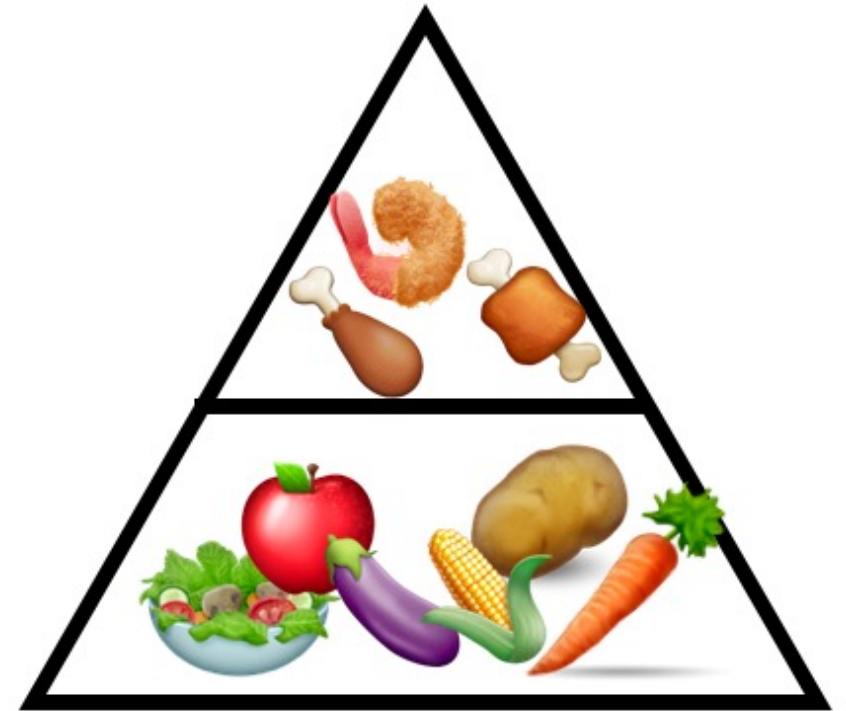


BASIC PCA ALGORITHM

- Start with a zero-centered $m \times n$ data matrix **X**
- Compute covariance matrix
- Find eigenvectors of covariance matrix
- PCs: k eigenvectors with highest eigenvalues (variance)

EXAMPLE: FOOD NUTRITION

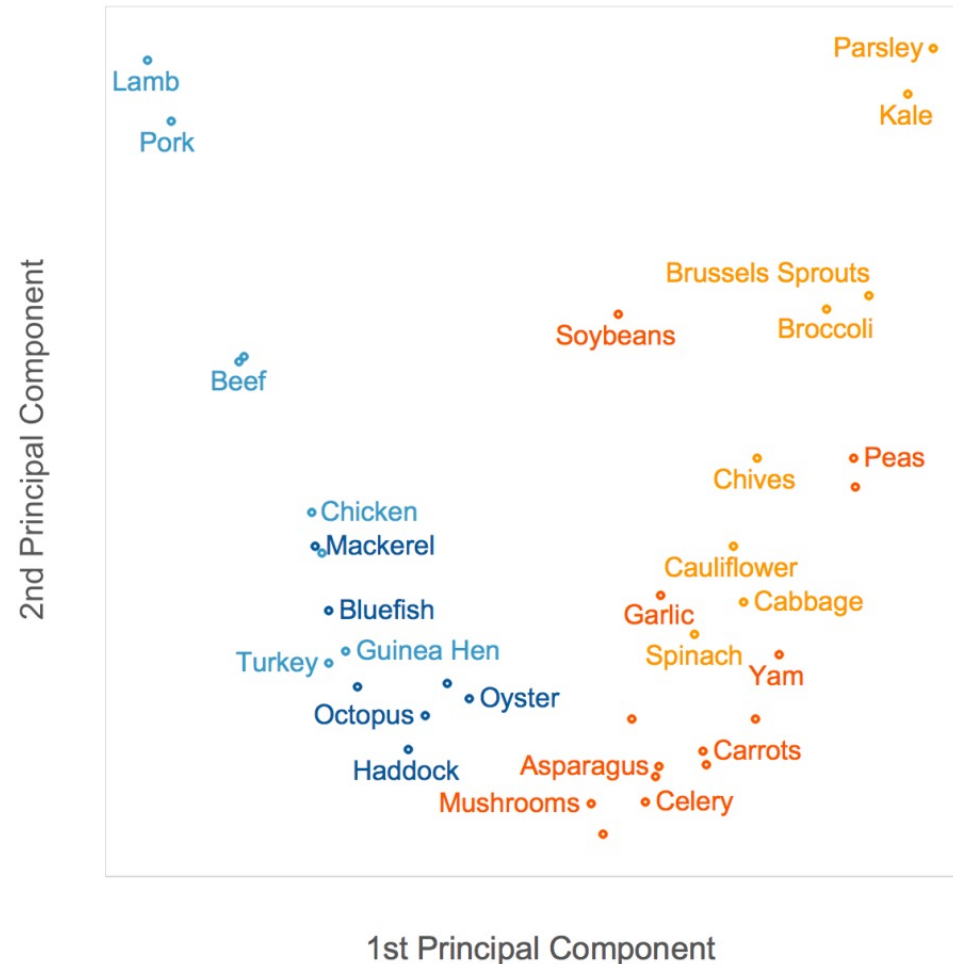
- What is the best way to differentiate food items?
- Vitamin content
- Protein levels
- Fat
- Fiber



<https://algobeans.com/2016/06/15/principal-component-analysis-tutorial/>

EXAMPLE: PCA

	PC1	PC2	PC3	PC4
Fat	-0.45	0.66	0.58	0.18
Protein	-0.55	0.21	-0.46	-0.67
Fiber	0.55	0.19	0.43	-0.69
Vitamin C	0.44	0.70	-0.52	0.22



<https://algobeans.com/2016/06/15/principal-component-analysis-tutorial/>

PCA: # OF PCS?

- How many components are sufficient to summarize the data?

PROPORTION VARIANCE EXPLAINED

- Total variance in data (assuming zero mean):

$$\text{Var}(\mathbf{X}) = \sum_{j=1}^p \frac{1}{n} \sum_{i=1}^n x_{ij}^2$$

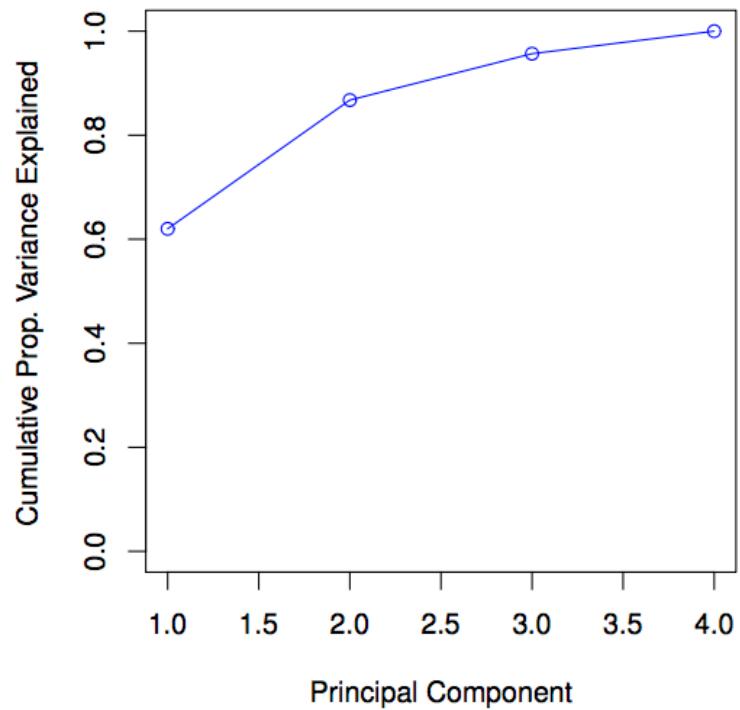
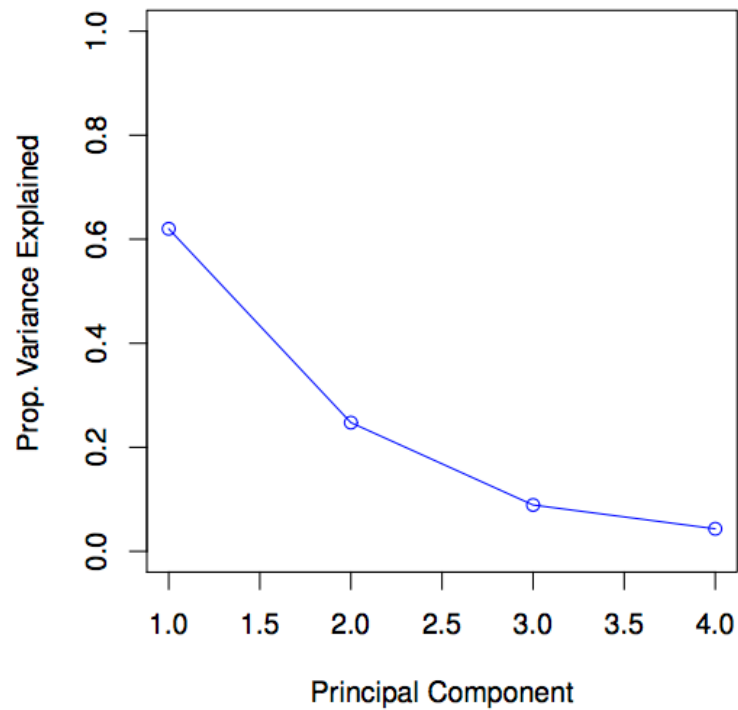
- Variance explained by the m^{th} component:

$$\text{Var}(\mathbf{W}_m) = \frac{1}{n} \sum_{i=1}^n w_{im}^2$$

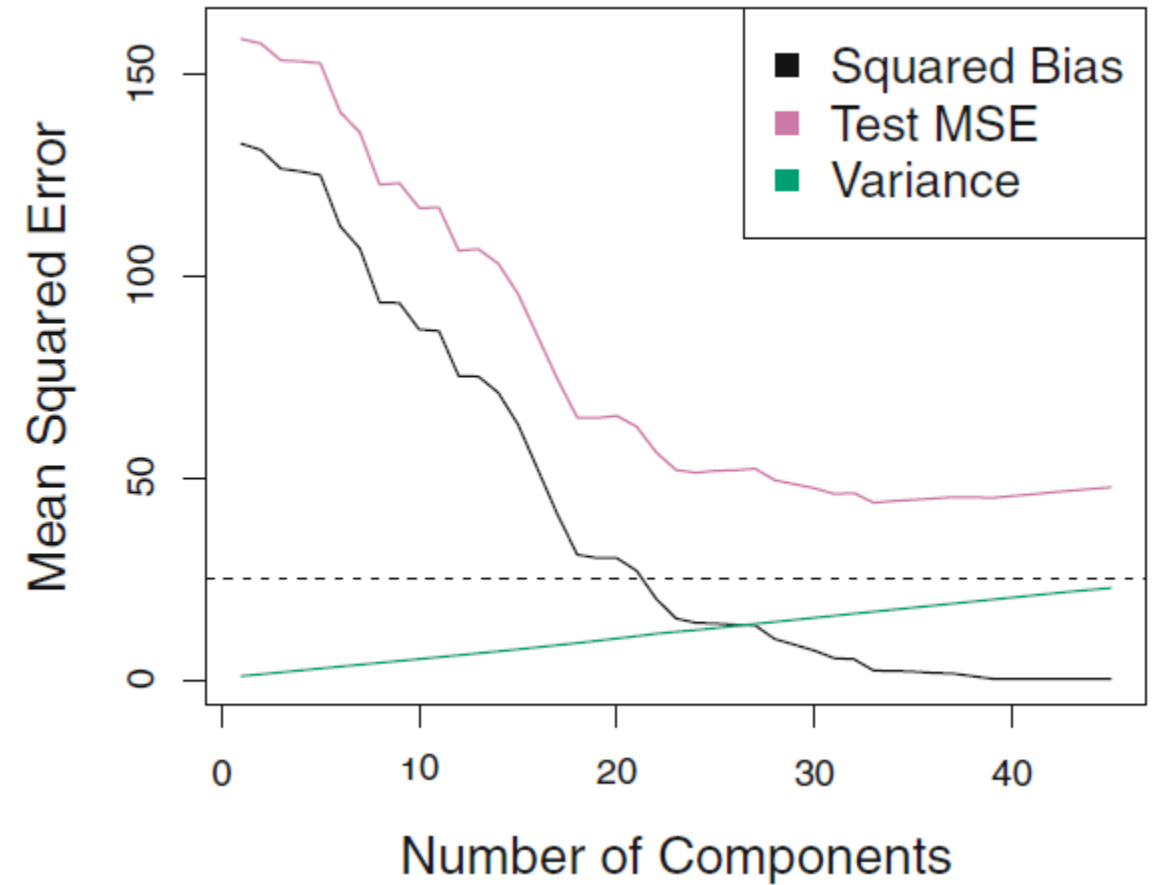
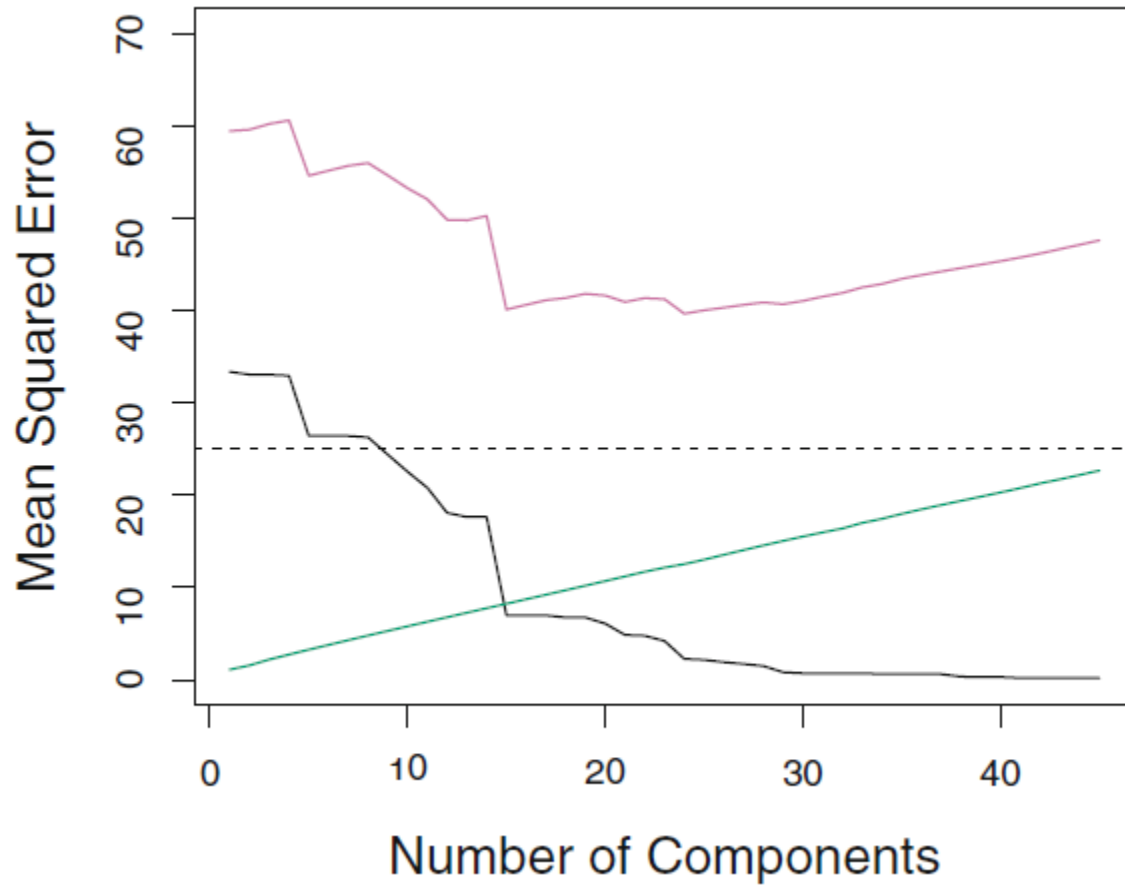
- Proportion of variance explained by m^{th} component :

$$\text{PVE}_m = \frac{\text{Var}(\mathbf{W}_m)}{\text{Var}(\mathbf{X})}$$

PCA: SCREE PLOT (UNSUPERVISED)

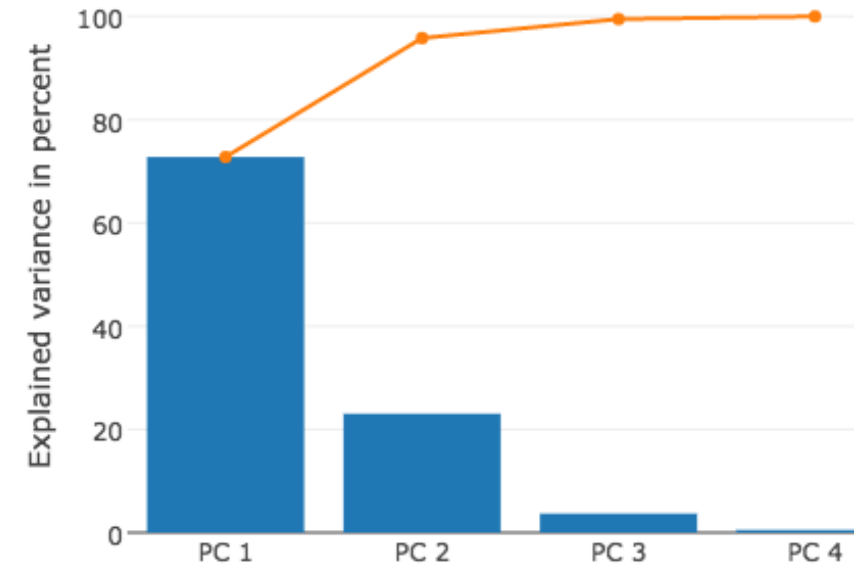


USE PCA FOR SUPERVISED LEARNING



PCA: INTERPRETATION

- If variances of PCs drop off quickly, then X is highly collinear
- Reduce dimensionality of data by keeping only the PCs with highest variance



<https://plot.ly/ipython-notebooks/principal-component-analysis/>

PCA: SKLEARN

```
from sklearn.decomposition import PCA as sklearnPCA
sklearn_pca = sklearnPCA(n_components=2)
Y_sklearn = sklearn_pca.fit_transform(X_std)
```

HOMEWORK #5

- Due 11/16 @ 11:59 PM ET on Gradescope
- 2 questions
 - PCA
 - Almost Random Forest



DEMO: PCA-EXAMPLE.IPYNB

[HTTPS://COLAB.RESEARCH.GOOGLE.COM/DRIVE/12VDO-jD0PVFLVZZLI2FT50DYESj8V6WWW](https://colab.research.google.com/drive/12VDO-jD0PVFLVZZLI2FT50DYESj8V6WWW)