

# SPOTLIGHT GRADES

	Min	Max	Mean	Median
Class average	8.79	9.55	9.32	9.26
Average (class average, instructor score)	8.74	9.68	9.46	9.35

Excellent job!

# MIDTERM

- 11/8 Wednesday in class 1-2:15pm
- Open book/notes, no digital device
- Refer to midterm review and sample questions
- All ML algorithms: basic algorithm (intermediate results of each step); decision boundary; impact of key parameters; distance metrics; regularization effect; bias and variance tradeoff
- Data preprocessing; model assessment and selection strategies

# TYPES OF UNSUPERVISED LEARNING

## Clustering

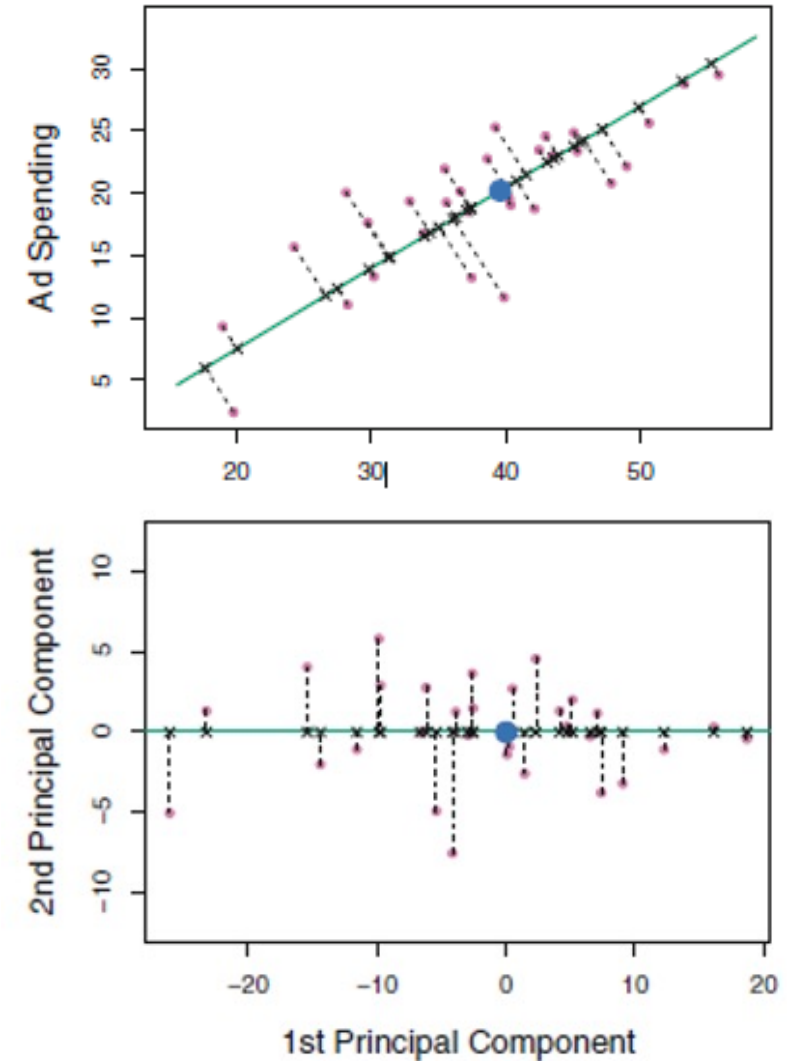
identify unknown structure in the data

## Dimensionality Reduction

use structural characteristics to simplify data

# REVIEW: PRINCIPAL COMPONENT ANALYSIS (PCA)

- First principal component
  - Yields the highest variance of the projection
- Second principal component
  - Orthogonal to the first principal component
  - And has largest variance
- In general, we can construct up to  $p$  principal components ( $p$  features)



# REVIEW: PCA: PROBLEM FORMULATION

- Given a feature matrix  $\mathbf{X}$  with  $n$  data points, find  $\mathbf{W}$  such that  $\|\mathbf{W}\|_2 = 1$  and the  $\text{Var}(\mathbf{XW})$  is maximized and  $\mathbf{W}$  consists of orthonormal vectors

$$\begin{aligned}\text{Var}(\mathbf{XW}) &= \frac{1}{N} (\mathbf{W}^\top (\mathbf{X} - \mu_{\mathbf{X}})^\top (\mathbf{X} - \mu_{\mathbf{X}}) \mathbf{W}) \\ &= \mathbf{W}^\top \Sigma_{\mathbf{X}} \mathbf{W}\end{aligned}$$

Sample covariance matrix

- The solution is the Eigenvectors of the covariance matrix
  - PCs:  $k$  eigenvectors with highest eigenvalues (variance)

# PCA: # OF PCS?

- How many components are sufficient to summarize the data?

# PROPORTION VARIANCE EXPLAINED

- Total variance in data (assuming zero mean):

$$\text{Var}(\mathbf{X}) = \sum_{j=1}^p \frac{1}{n} \sum_{i=1}^n x_{ij}^2$$



- Variance explained by the  $m^{\text{th}}$  component:

$$\text{Var}(\mathbf{W}_m) = \frac{1}{n} \sum_{i=1}^n w_{im}^2$$

- Proportion of variance explained by  $m^{\text{th}}$  component :

$$\text{PVE}_m = \frac{\text{Var}(\mathbf{W}_m)}{\text{Var}(\mathbf{X})}$$


# PCA: SCREE PLOT (UNSUPERVISED)







# PCA: INTERPRETATION

- If variances of PCs drop off quickly, then  $X$  is highly collinear
- Reduce dimensionality of data by keeping only the PCs with highest variance



<https://plot.ly/ipython-notebooks/principal-component-analysis/>

# USE PCA FOR SUPERVISED LEARNING



# PCA: SKLEARN

```
from sklearn.decomposition import PCA as sklearnPCA
sklearn_pca = sklearnPCA(n_components=2)
Y_sklearn = sklearn_pca.fit_transform(X_std)
```

# REMINDER: HOMEWORK #5

- Due 11/16 @ 11:59 PM ET on Gradescope
- 2 questions
  - PCA
  - Almost Random Forest



DEMO: PCA-EXAMPLE.IPYNB

[HTTPS://COLAB.RESEARCH.GOOGLE.COM/DRIVE/12VDO-jD0PVFLVZZLI2FT50DYESj8V6WWW](https://colab.research.google.com/drive/12VDO-jD0PVFLVZZLI2FT50DYESj8V6WWW)

# TYPES OF UNSUPERVISED LEARNING

## Clustering


identify unknown structure in the data

## Dimensionality Reduction

use structural characteristics to simplify data

# WHAT IS CLUSTER ANALYSIS?

- Finding groups of objects (clusters)
- Objects **similar** to one another in the same group
- Objects **different** from the objects in other groups
- Unsupervised learning: no predefined classes



How to define similarity?

# APPLICATIONS OF CLUSTER ANALYSIS

- As a stand-alone tool to get insight into data distribution
  - Cluster into groups – automatic classification
  - Finding k-nearest neighbors
  - Outlier detection
- As a preprocessing step for other algorithms
  - Data cleaning: missing data, noisy data; Data reduction; Data discretization
  - Build supervised models for each cluster




# CLUSTERING APPROACHES

- Partitioning approach:
  - Construct various partitions and then evaluate them by some “goodness” criterion
  - Typical methods: **k-means**, k-medoids
- Hierarchical approach:
  - Create a hierarchical decomposition of the objects
  - Typical methods: Diana, Agnes
- Density-based approach:
  - Based on connectivity and density functions (vs. distance)
  - Typical methods: DBSACN
- Others




GROUP ACTIVITY

Which is better and why?  
What is good clustering criteria?



# QUALITY: WHAT IS GOOD CLUSTERING?


- Homogeneity - high intra-class similarity
- Separation - low inter-class similarity



# PARTITIONING ALGORITHMS: BASIC CONCEPT

- Partitioning method: Construct a partition of  $n$  objects (into  $k$  clusters), s.t. intra-cluster similarity maximized and inter-cluster similarity minimized
- One objective: minimize the sum of squared distance from **cluster centroid** (intra-class distance)

$$\sum_{i=1}^k \sum_{p \in C_i} (p - m_i)^2$$



- How to find optimal partition?

# NUMBER OF PARTITIONINGS

- Stirling partition number – number of ways to partition  $n$  objects into  $k$  non-empty subsets

$$\left\{ \begin{matrix} n+1 \\ k \end{matrix} \right\} = k \left\{ \begin{matrix} n \\ k \end{matrix} \right\} + \left\{ \begin{matrix} n \\ k-1 \end{matrix} \right\}$$


( $n=5$ ,  $k=1, 2, 3, 4, 5$ ): 1, 15, 25, 10, 1

( $n=10$ ,  $k=1, 2, 3, 4, 5, \dots$ ): 1, 511, 9330, 34105, 42525, ...

- Bell numbers – number of ways to partition  $n$  objects

$$B_n = \sum_{k=0}^n \left\{ \begin{matrix} n \\ k \end{matrix} \right\}.$$

( $n=0, 1, 2, 3, 4, 5, \dots$ ): 1, 1, 2, 5, 15, 52, 203, 877, 4140, 21147, 115975, 678570, 4213597, 27644437, 190899322, 1382958545, 10480142147, 82864869804, 682076806159, 5832742205057, ...




# *K-MEANS* CLUSTERING: LLOYD ALGORITHM

- Lloyd'57, MacQueen'67
- Heuristic EM-style algorithm for the partitioning problem
- Each cluster is represented by the center of the cluster

# K-MEANS CLUSTERING: LLOYD ALGORITHM

(LLOYD'57, MACQUEEN'67)

- If we know the centroid of each cluster, how do we cluster the points?






# K-MEANS CLUSTERING: LLOYD ALGORITHM

(LLOYD'57, MACQUEEN'67)


- If we know the centroid of each cluster, how do we cluster the points?
- If we know the clusters, how do we compute centroid?



# K-MEANS CLUSTERING: LLOYD ALGORITHM

(LLOYD'57, MACQUEEN'67)

- If we know the centroid of each cluster, how do we cluster the points?
- If we know the clusters, how do we compute centroid?





How do we get started?


# K-MEANS CLUSTERING: LLOYD ALGORITHM

(LLOYD'57, MACQUEEN'67)







- Initialization: Given  $k$ , randomly choose  $k$  initial cluster centers
- **Assignment**: Partition objects into  $k$  nonempty subsets by assigning each object to the cluster with the **nearest** centroid
- **Update**: update centroid, i.e. *mean point* of the cluster
- Go back to Step 2 and repeat, stop when no more new assignment and centroids do not change



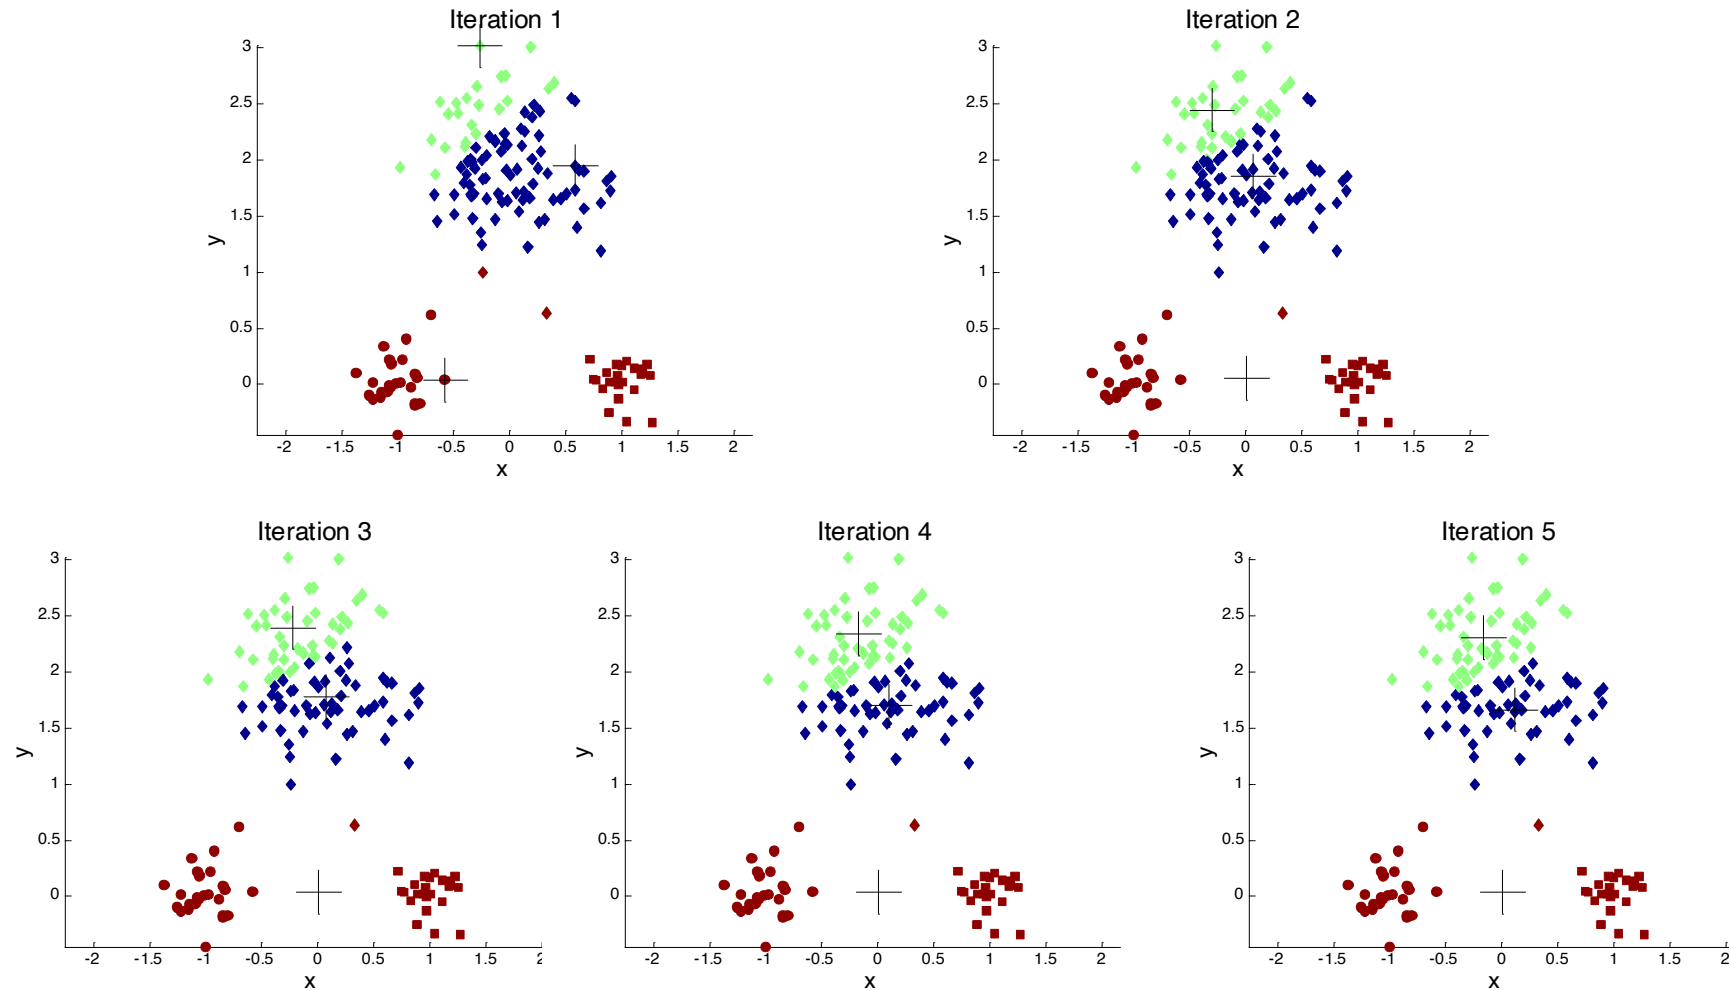
# THE *K*-MEANS CLUSTERING METHOD



# IMPORTANCE OF CHOOSING INITIAL CENTROIDS – CASE I



# IMPORTANCE OF CHOOSING INITIAL CENTROIDS – CASE 2



How do we avoid bad initial cases?

# K-MEANS CLUSTERING – DETAILS

- Initial centroids are often chosen randomly
  - Example: Pick one point at random, then  $k-1$  other points, each as far away as possible from the previous points
- Run multiple times
- ‘Nearest’ is measured by Euclidean distance, cosine similarity, correlation, etc.
- Most of the convergence happens in the first few iterations.
  - Often the stopping condition is changed to ‘Until relatively few points change clusters’
- What’s the complexity? (Minibatch k-means further reduces complexity)  
 $n$  is # objects,  $k$  is # clusters, and  $t$  is # iterations.

# CLUSTERING EVALUATION

- SSE (sum of squared error)

$$\sum_{i=1}^k \sum_{p \in C_i} (p - m_i)^2$$

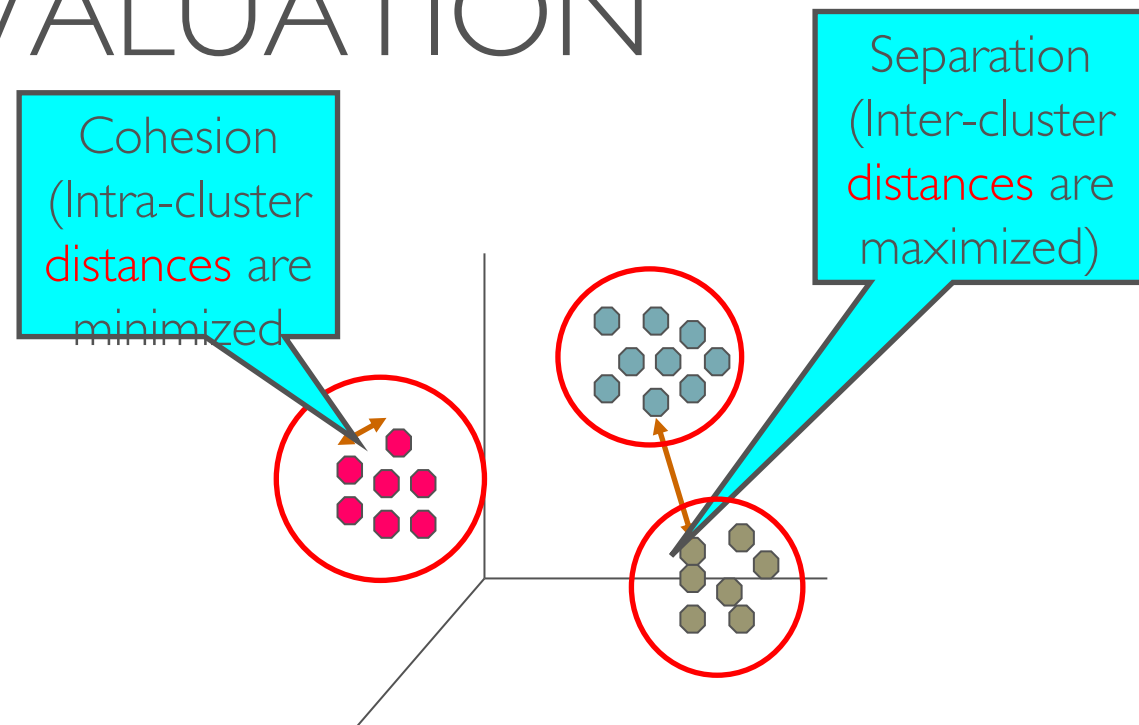
- Silhouette coefficient  $[-1, 1]$

$a(i)$ : The mean distance between a sample  $i$  and all other points in the same cluster

$b(i)$ : The mean distance between a sample  $i$  and all other points in the next nearest cluster.

$$a(i) = \frac{1}{|C_I| - 1} \sum_{j \in C_I, i \neq j} d(i, j) \quad b(i) = \min_{J \neq I} \frac{1}{|C_J|} \sum_{j \in C_J} d(i, j)$$

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}, \text{ if } |C_I| > 1$$






# CLUSTERING EVALUATION - SUPERVISED

- Compare clusters with “ground truth” clusters

- Entropy/purity based
- Precision and recall based
- Similarity-based



► **Figure 16.1** Purity as an external evaluation criterion for cluster quality. Majority class and number of members of the majority class for the three clusters are: x, 5 (cluster 1); o, 4 (cluster 2); and  $\diamond$ , 3 (cluster 3). Purity is  $(1/17) \times (5 + 4 + 3) \approx 0.71$ .

# PRECISION AND RECALL BASED

- BCubed Precision and recall – average precision and recall of all objects
  - Precision of an object: proportion of objects in the same cluster belong to the same category
  - Recall of an object: proportion of objects of the same category are assigned to the same cluster




Figure 10: Example of computing the BCubed precision and recall for one item


# SIMILARITY-BASED MEASURES

Given a reference clustering T and clustering S

- $f_{00}$ : number of pair of points belonging to different clusters in both T and S
- $f_{01}$ : number of pair of points belonging to different cluster in T but same cluster in S
- $f_{10}$ : number of pair of points belonging to same cluster in T but different cluster in S
- $f_{11}$ : number of pair of points belonging to same clusters in both T and S

$$Rand = \frac{f_{00} + f_{11}}{f_{00} + f_{01} + f_{10} + f_{11}}$$

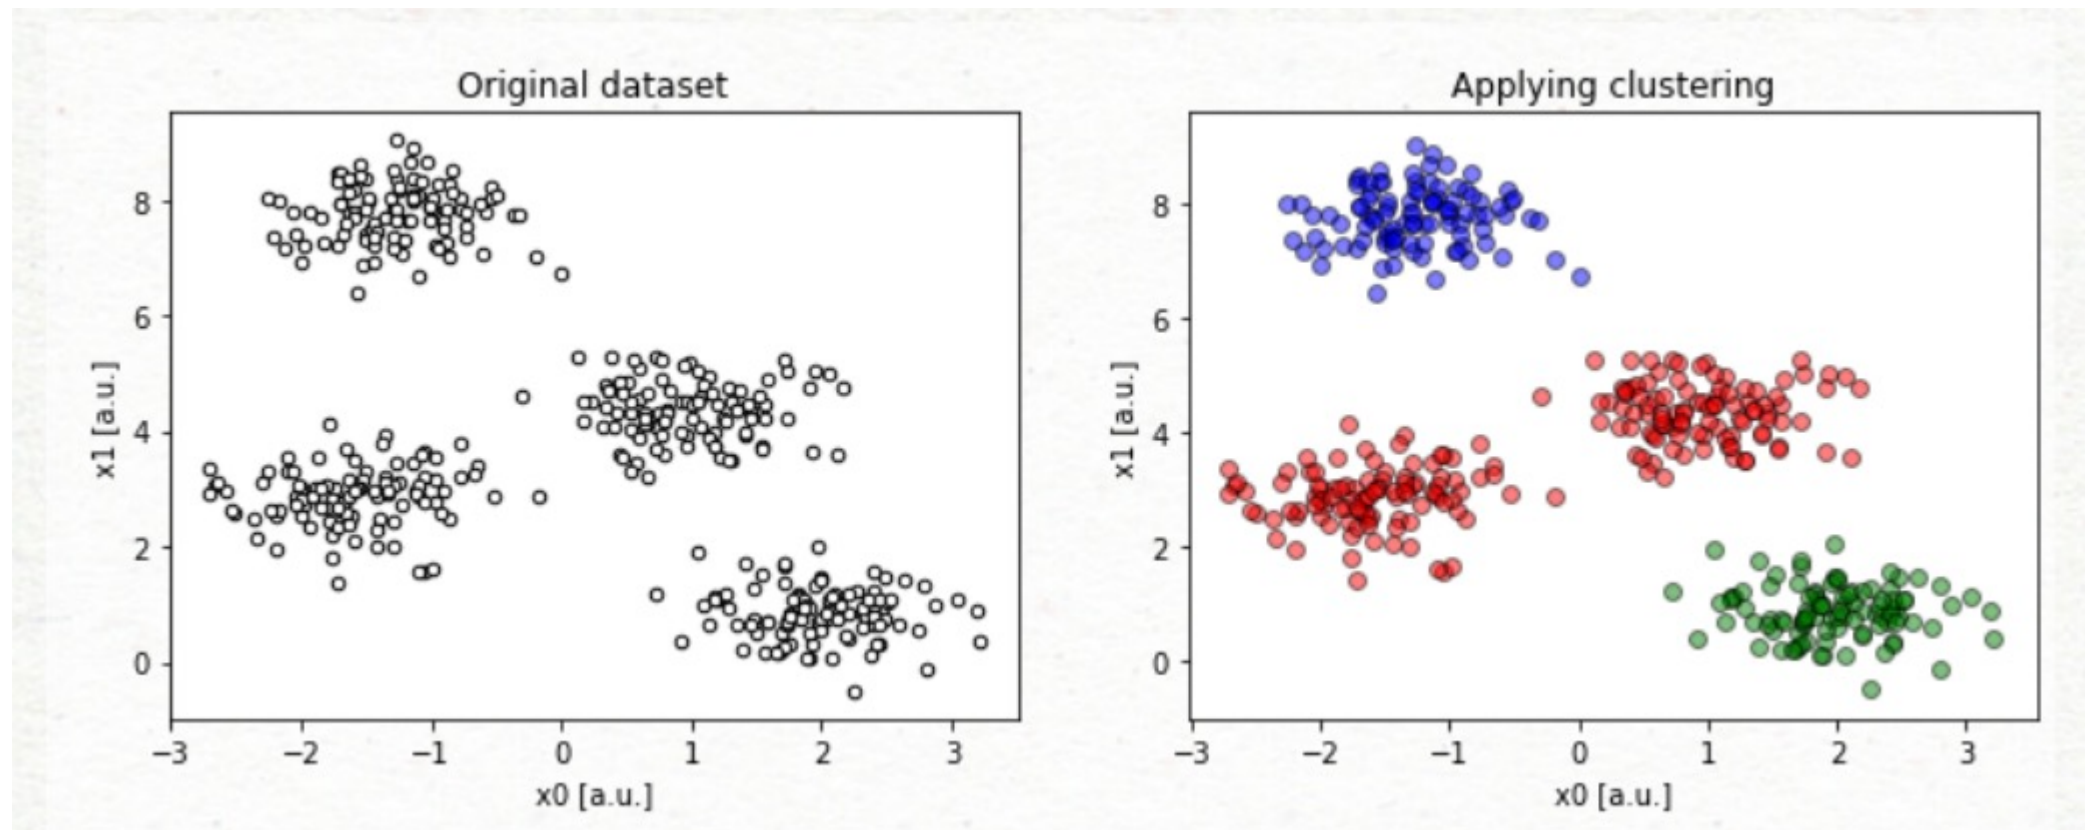
$$Jaccard = \frac{f_{11}}{f_{01} + f_{10} + f_{11}}$$




# K-MEANS CLUSTERING: # OF K?

- How to choose the number of clusters K?

# K: TOO FEW




# K: TOO MANY





# GETTING THE $K$ RIGHT – ELBOW METHOD

- Try different  $k$ , looking at the change in clustering criterion (e.g. average distance to centroid, SSE, or clustering coefficient) as  $k$  increases
- Average distance typically falls rapidly until right  $k$ , then changes little




# CHOOSING K - ELBOW METHOD




SSE of clusters found using K-means



# LIMITATIONS OF K-MEANS: DIFFERING SIZES




Original Points




K-means (3 Clusters)

# LIMITATIONS OF K-MEANS: NON-CONVEX SHAPES




Original Points




K-means (2 Clusters)

# OVERCOMING K-MEANS LIMITATIONS




Original Points




K-means Clusters

# OVERCOMING K-MEANS LIMITATIONS



Original Points



K-means Clusters

Or use hierarchical or density based clustering

# K-MEANS CLUSTERING: SKLEARN

```
from sklearn.cluster import KMeans  
kmeans = KMeans(n_clusters=8, n_init=10)  
kmeans.fit(X)  
print(kmeans.labels_)
```

- n\_clusters: number of clusters
- n\_init: number of times k-means is run